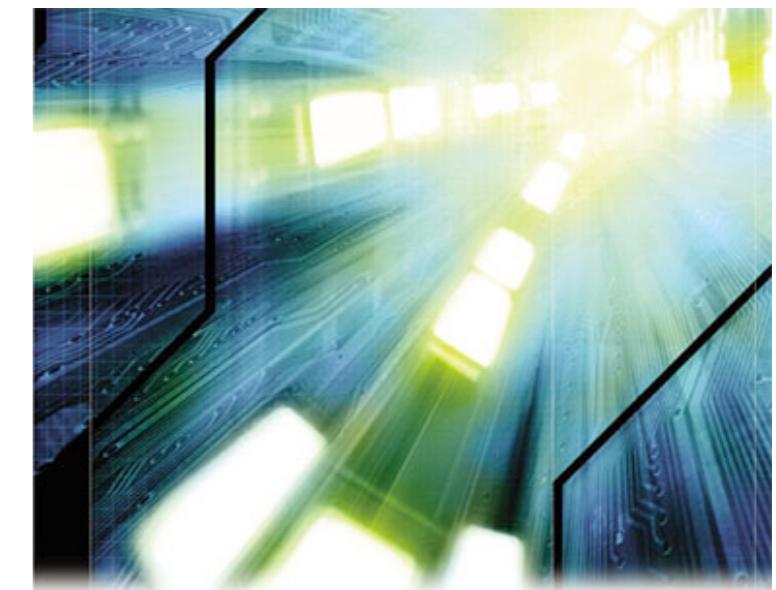


System Models

Nicola Dragoni
Embedded Systems Engineering
DTU Informatics

-
- 2.1 Introduction
 - 2.2 Architectural Models
 - 2.3 Fundamental Models



fourth edition

DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg



Architectural vs Fundamental Models

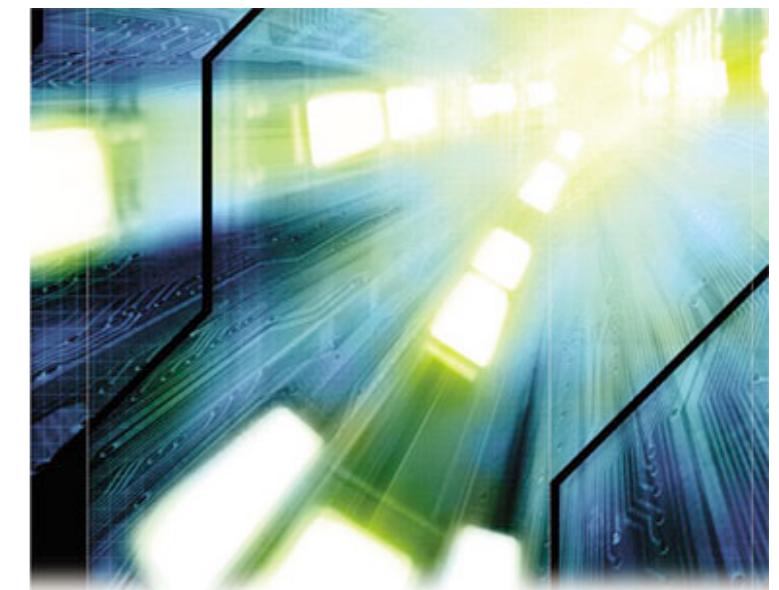
- Systems that are intended for use in real-world environments should be designed to function correctly in the widest possible range of circumstances and in the face of many possible difficulties and threats.
- An architectural model is concerned with the placement if its components and the relationships between them.
 - ▶ client-server systems
 - ▶ peer-to-peer systems

Architectural vs Fundamental Models

- Systems that are intended for use in real-world environments should be designed to function correctly in the widest possible range of circumstances and in the face of many possible difficulties and threats.
- An architectural model is concerned with the placement if its components and the relationships between them.
 - ▶ client-server systems
 - ▶ peer-to-peer systems
- Fundamental models are concerned with a more formal description of the properties that are common in all of the architectural models.

System Models

-
- 2.1 Introduction
 - 2.2 Architectural Models
 - 2.3 Fundamental Models



fourth edition

DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg



Architectural Models



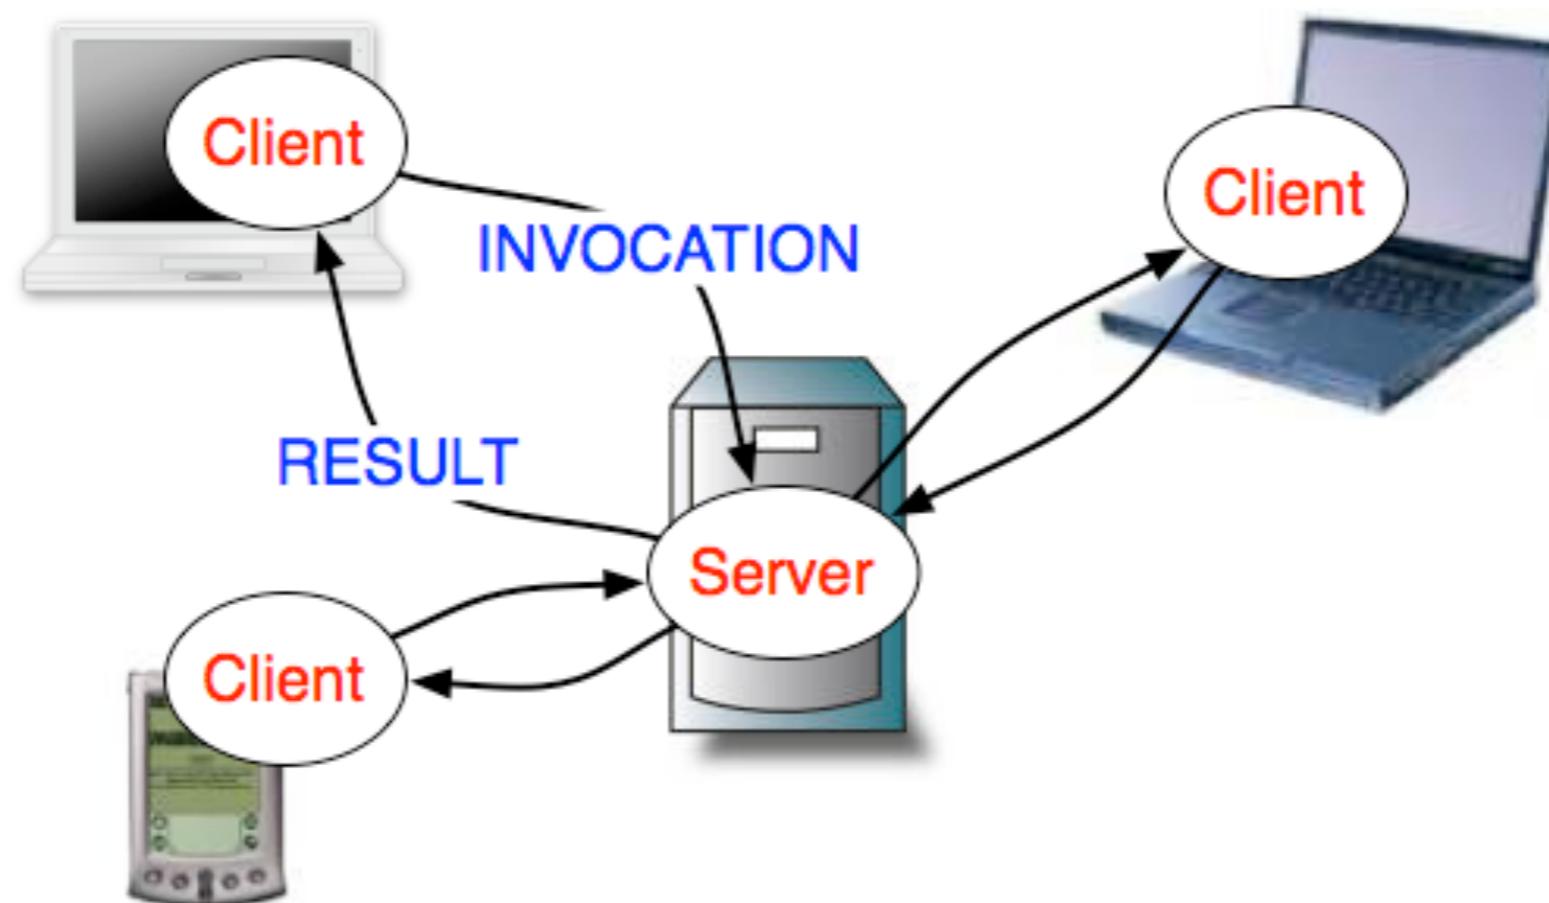
- The **architecture of a system** is its structure in terms of separately specified components.
- An **architectural model** of a distributed system:
 1. it **simplifies and abstracts the functions** of the **individual components**
 2. it considers:
 - ▶ the **placement of the components** across a network of computers (seeking to define useful patterns for the distribution of data and workload)
 - ▶ the **interrelationships between the components** (i.e., their functional roles and the patterns of communication between them).

Initial Classification

- Achieved by classifying processes as **server processes**, **client processes** and **peer processes** (processes that cooperate and communicate in a *symmetrical manner* to perform a task).
- This classification:
 - ▶ identifies the **responsibilities** of each component
 - ▶ helps to **assess** its **workload**
 - ▶ helps to **determine** the **impact of failures** in each component
- The results of this analysis can then be used to **specify** the **placement** of the processes in a manner that meets **performance** and **reliability goals** for the resulting system.

Architectural Model: Client-Server

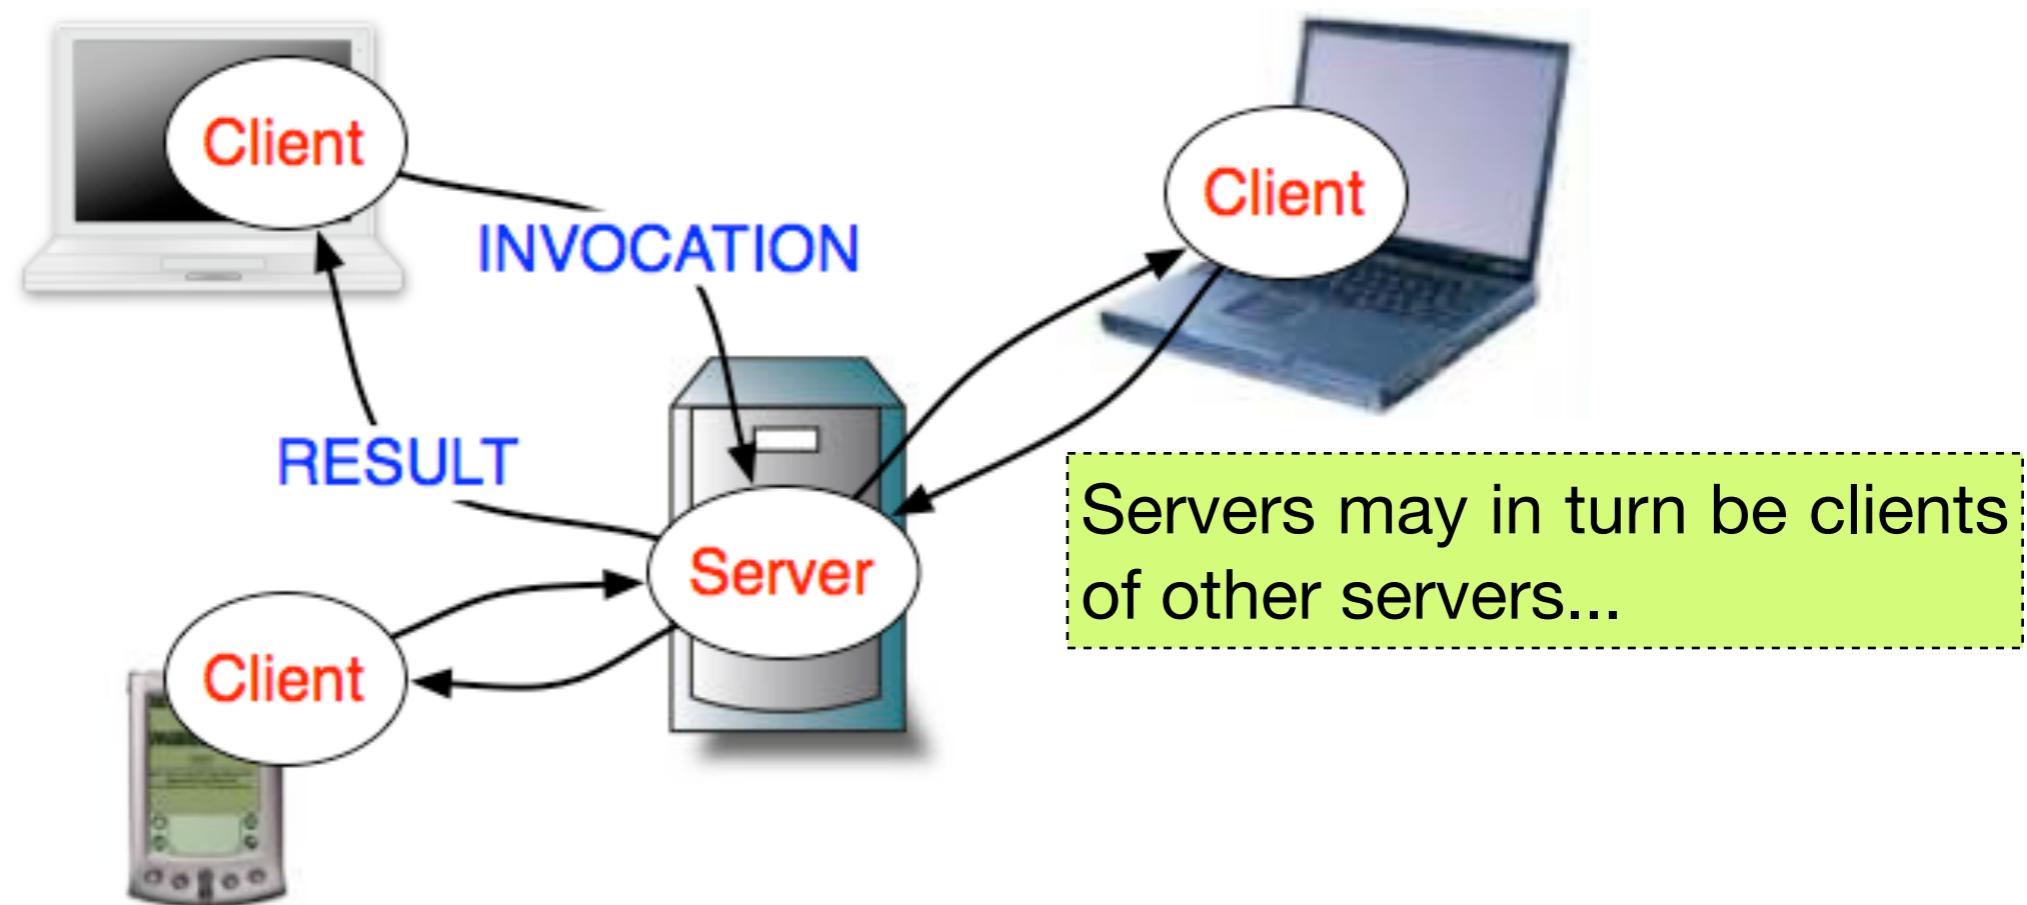
- Still the most widely employed architectural model.



Client processes interact with individual server processes in separate host computers in order to access the shared resources that they manage.

Architectural Model: Client-Server

- Still the most widely employed architectural model.



Client processes interact with individual server processes in separate host computers in order to access the shared resources that they manage.

On the Client-Server Role: Web Server Example

- Example 1: a **Web server** is often a **client** of a **local file server** that manages the files in which the web pages are stored.

On the Client-Server Role: Web Server Example

- Example 1: a **Web server** is often a **client** of a **local file server** that manages the files in which the web pages are stored.
- Example 2: **Web servers** and most **Internet services** are **clients** of the **DNS service** (which translates Internet Domain names to network addresses).

On the Client-Server Role: Web Server Example

- Example 1: a **Web server** is often a **client** of a **local file server** that manages the files in which the web pages are stored.
 - Example 2: **Web servers** and most **Internet services** are **clients** of the **DNS service** (which translates Internet Domain names to network addresses).
 - Example 3: **search engine**
 - ▶ **Server**: it responds to queries from browser clients
 - ▶ **Client**: it runs (in the background) programs called *web crawlers* that act as clients of other web serversThe diagram consists of two concentric circles. The inner circle contains the word "Google" in its signature blue, red, yellow, and green colors. The outer ring is divided into 12 segments, each containing a different search engine logo and name. Starting from the top and moving clockwise, the logos and names are: Teoma, looksmart, Lycos, HotBot, Go.com, Overture, AlltheWeb, MSN, Alexa, AOL, Netscape, and Yahoo!. Each segment is a different color, creating a vibrant, multi-colored wheel of search engines.



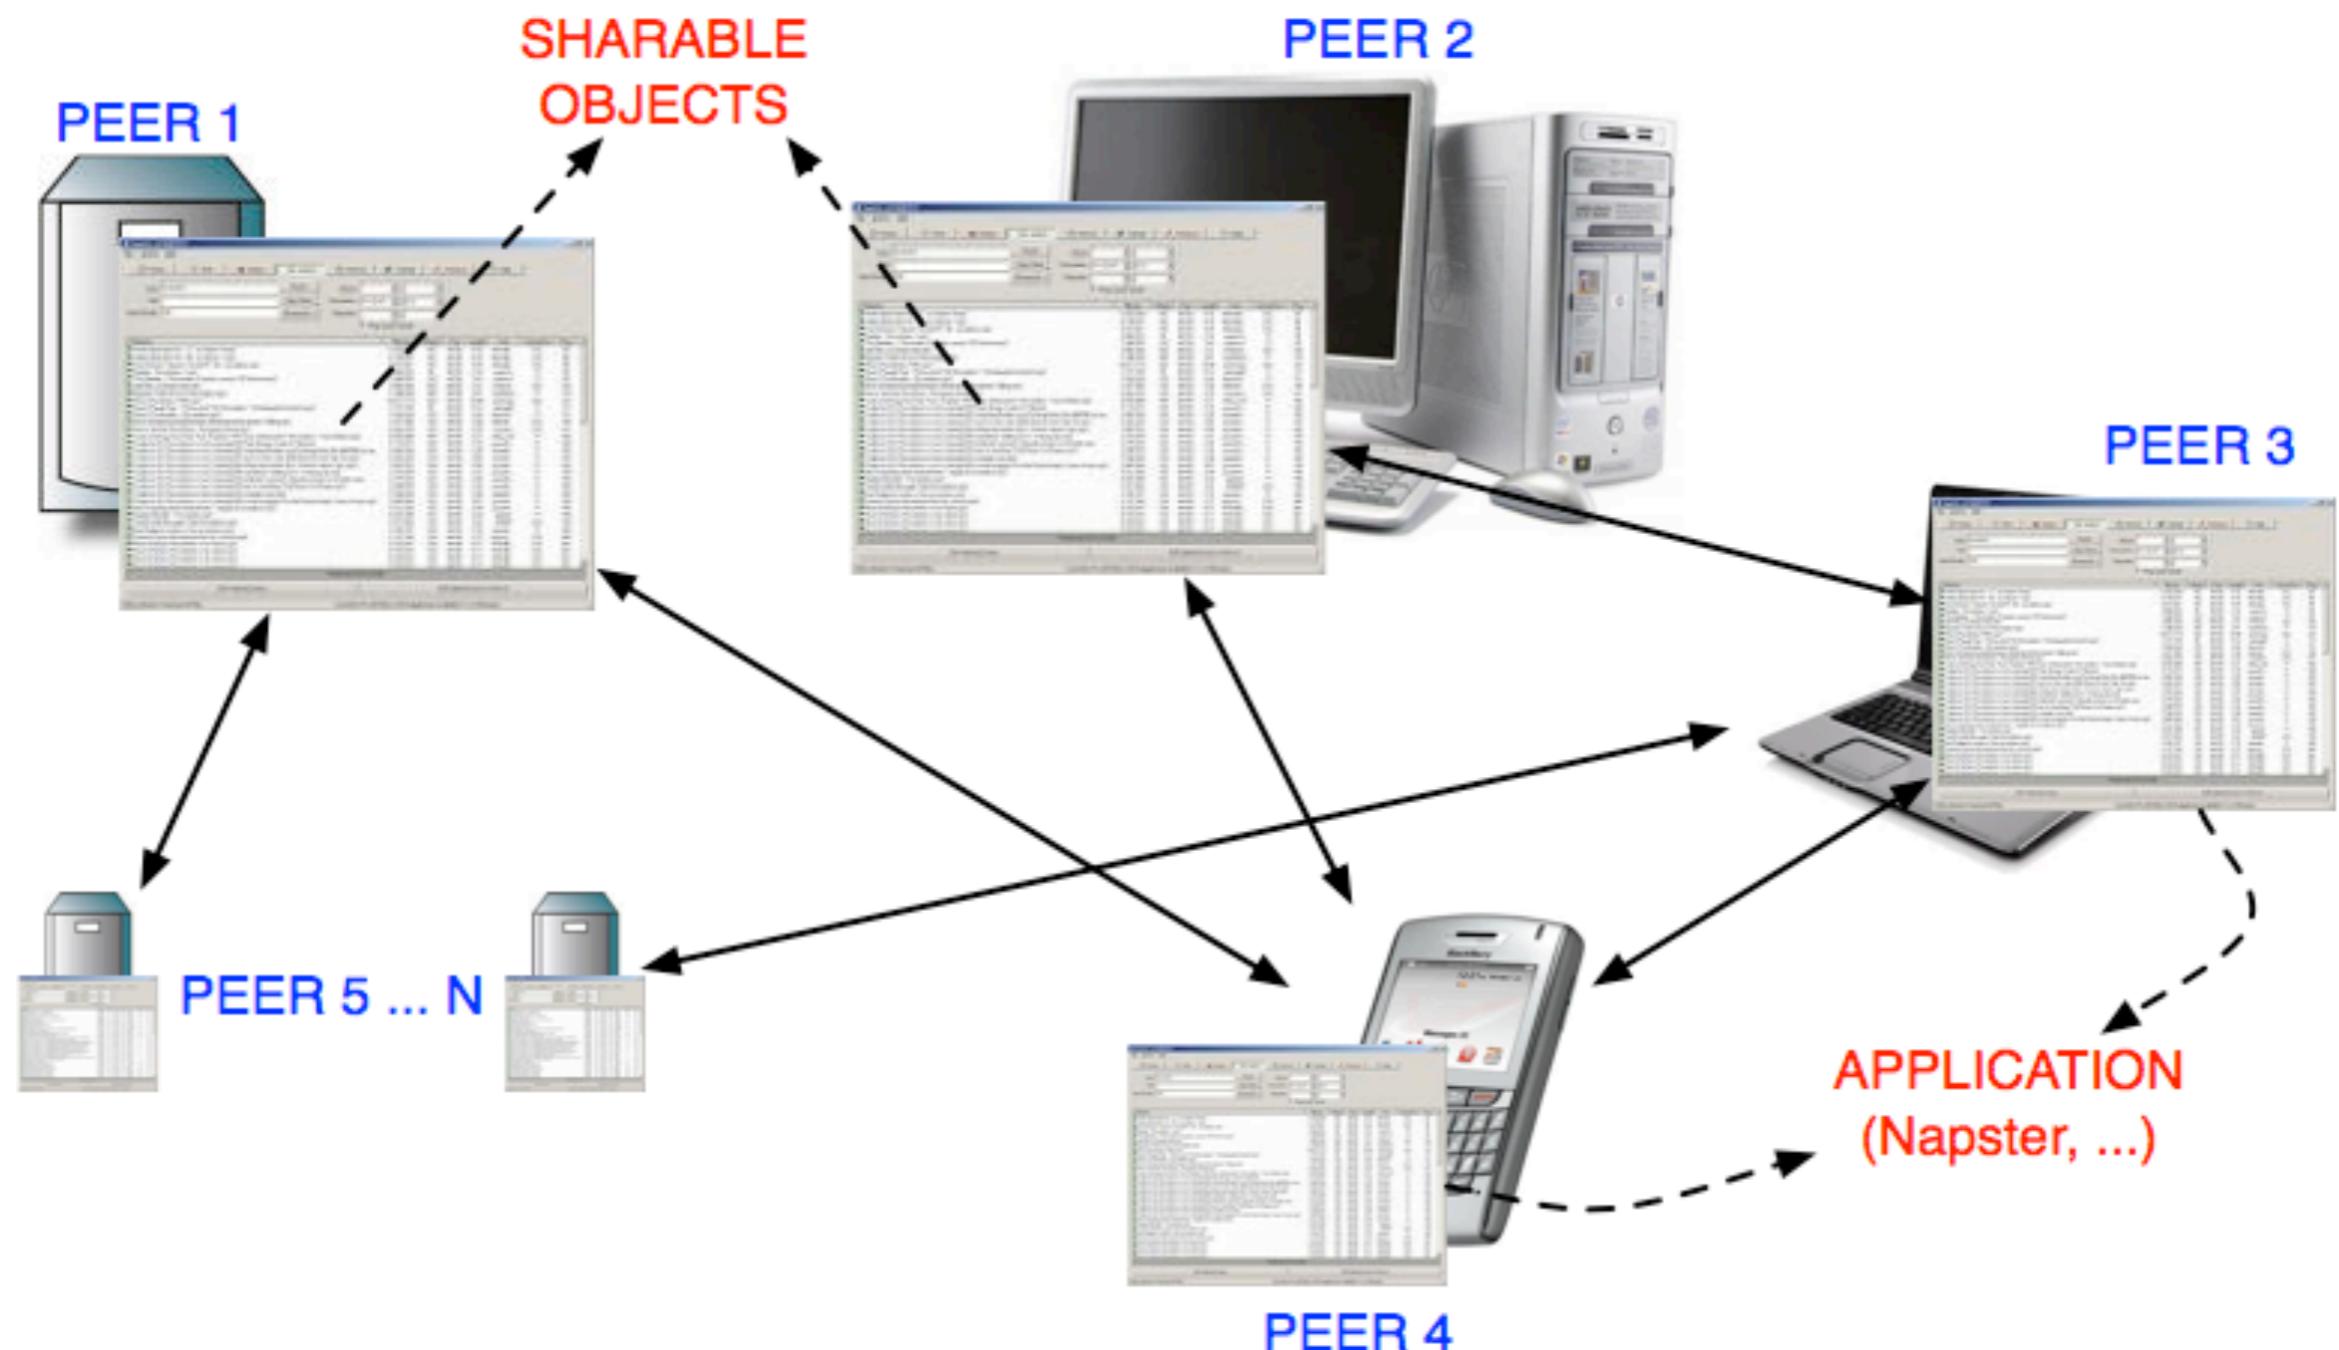
Architectural Model: Peer-to-Peer (P2P)

- All the processes involved in a task or activity play **similar roles**, interacting **cooperatively** as **peers** without any **distinction between client and server processes** or the computers that they run on.
- The **hardware capacity** and **operating system functionality** of today's desktop computers exceeds that of yesterday's servers.



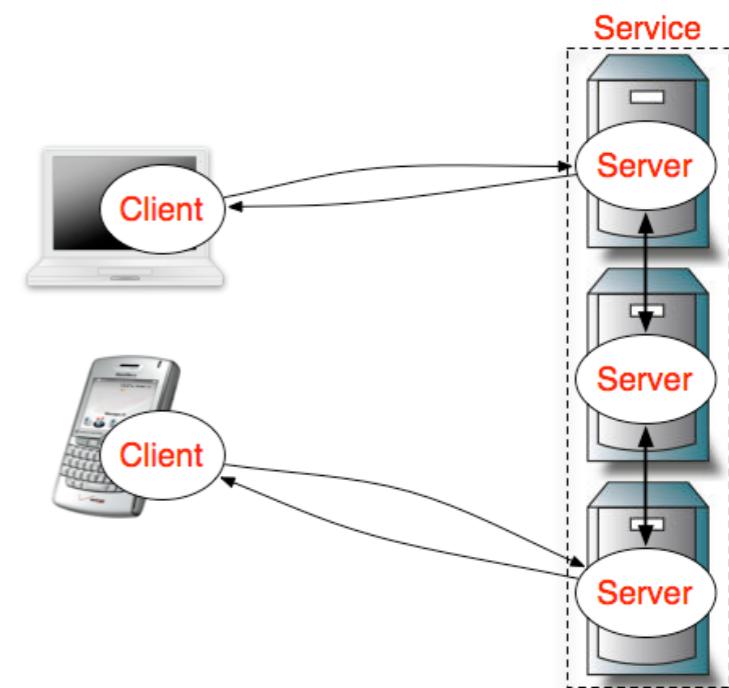
The **aim** of the P2P architecture is **to exploit the resources** (both data and hardware) in a large number of participating computers for the fulfillment of a given task or activity.

Distributed Application Based on a P2P Architecture

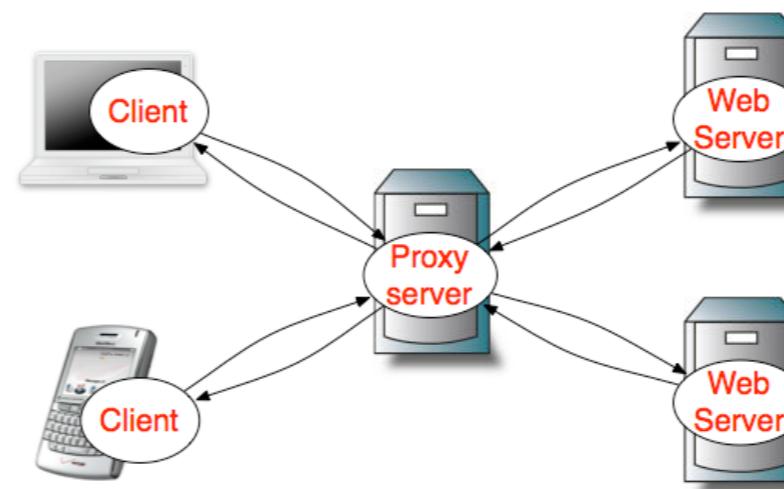


Variations of the Models

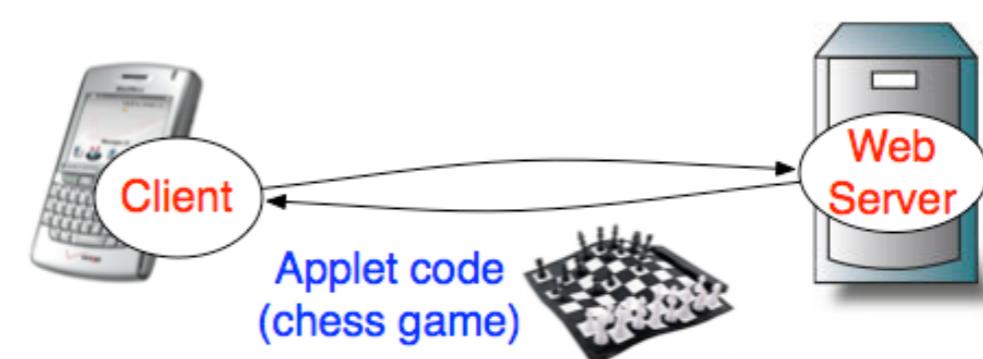
- Services provided by multiple servers



- Proxy server and caches



- Mobile code



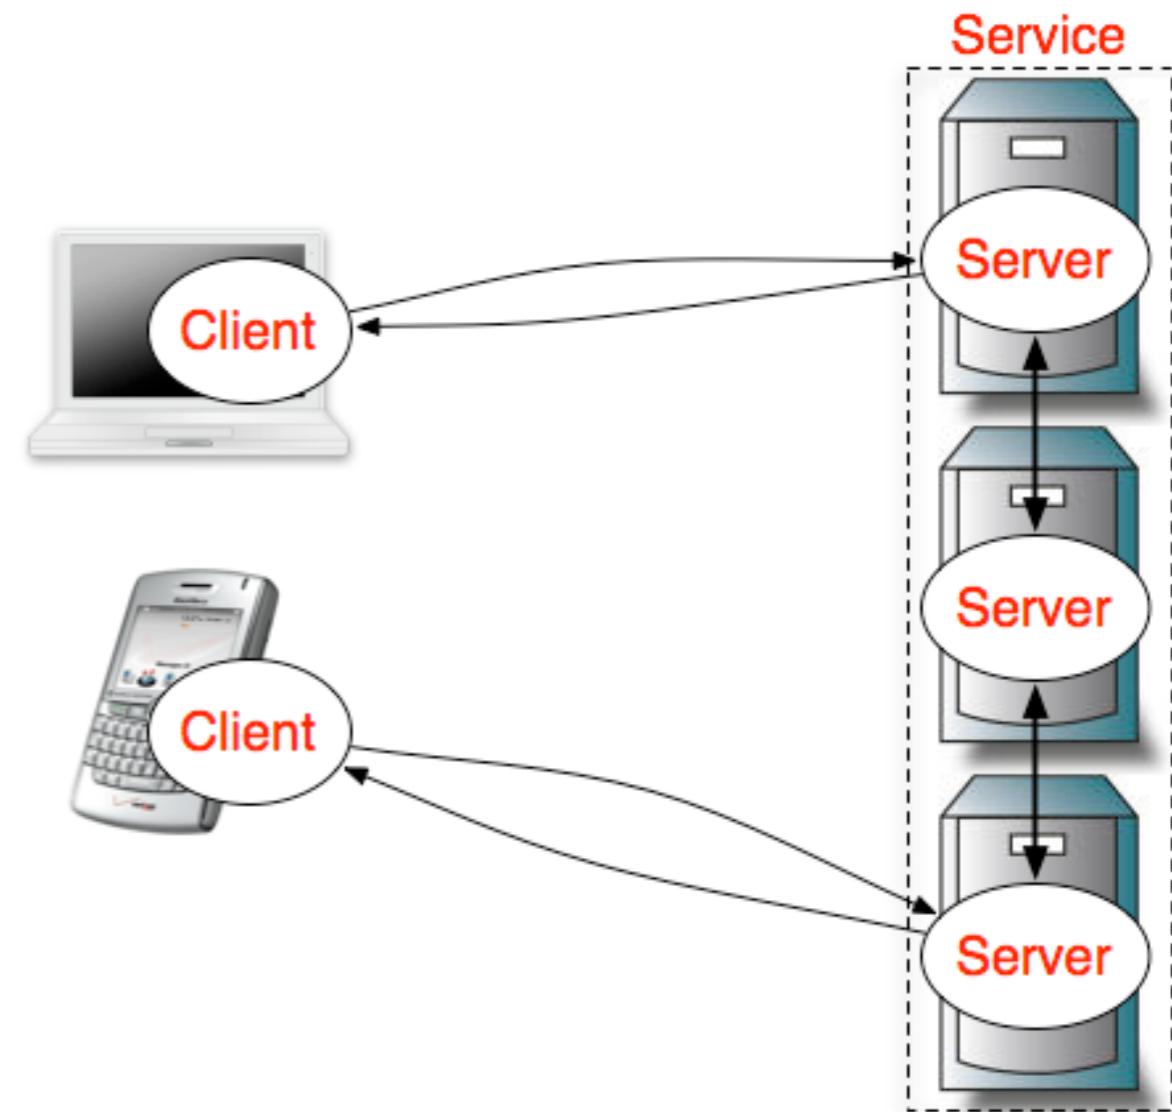
- ...

Variation: Services Provided by Multiple Servers

- Services may be implemented as several server processes in separate host computers interacting as necessary to provide a service to client processes.

- The servers may:

- 1) partition the set of objects on which the service is based and distributed them between themselves (e.g. Web servers)
- 2) they may maintain replicated copies of them on several hosts (e.g. SUN Network Information Service (NIS)).



Variation: Proxy Servers and Caches

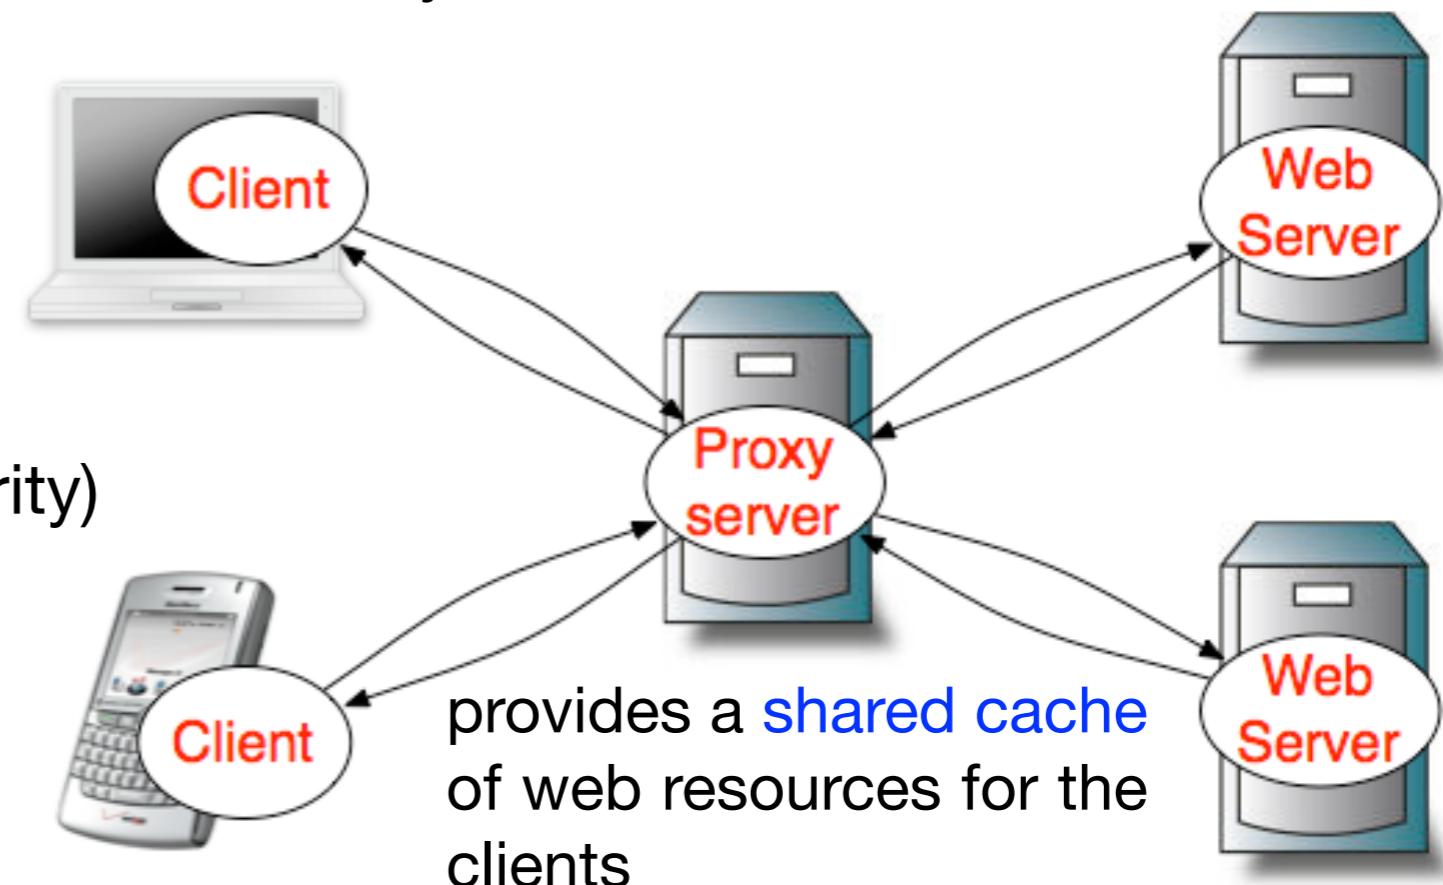
- A **cache** is a store of recently used data objects that is closer than the objects themselves.
- **Example 1:** Web browsers maintain a cache of recently visited pages and other web resources in the client's local file system.

Variation: Proxy Servers and Caches

- A **cache** is a store of recently used data objects that is closer than the objects themselves.
- Example 1: Web browsers maintain a cache of recently visited pages and other web resources in the client's local file system.
- Example 2: Web proxy server

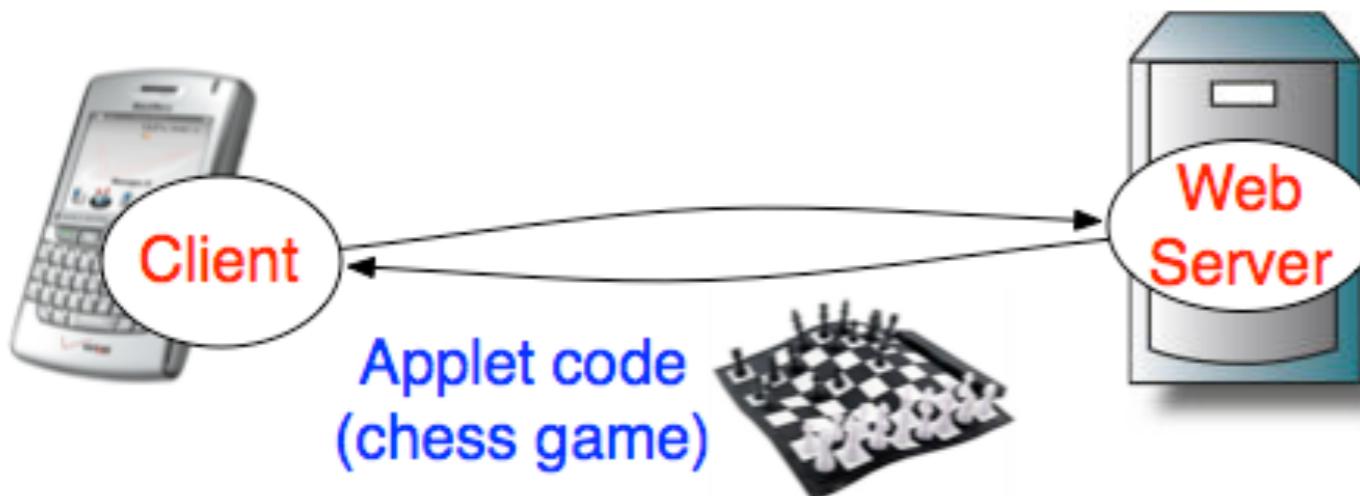
Purpose:

1. To keep machines behind it **anonymous** (mainly for security)
2. To **speed up** access to a resource (via caching)



Variation: Mobile Code

A) Client request results in the downloading of applet code



B) Client interacts with the applet



An **advantage** of running the downloaded code *locally* is that it can give good interactive response since it does not suffer from the delays or variability of bandwidth associated with network communication.

System Models

-
- 2.1 Introduction
 - 2.2 Architectural Models
 - 2.3 Fundamental Models



fourth edition

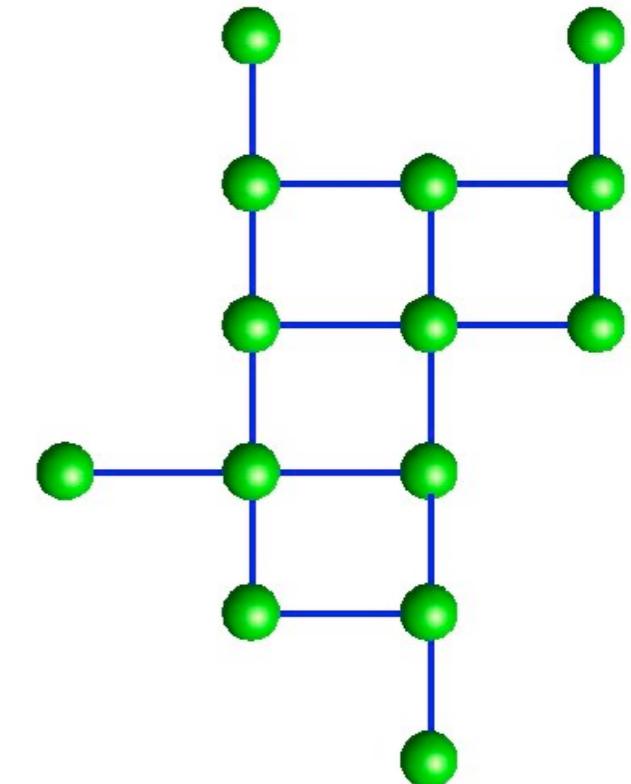
DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg



Fundamental (Abstract) Models

- The previous, quite different, models of systems share some **fundamental properties**.
 - ▶ For instance, all of them are composed of **processes** that **communicate** with one another by **sending messages** over a **computer network**.
- **Fundamental models** are concerned with a more formal description of the **properties** that are **common** in all the architectural models.



A **model** contains only the **essential ingredients** that we need to consider in order **to understand and reason about** some aspects of a system's behaviour.

Three Fundamental Models

- **Interaction model:** computation occurs within processes that interact by passing messages, resulting in **communication** (i.e., information flow) and **coordination** (synchronization and ordering of activities) between processes.
- **Failure model:** the correct operation of a distributed system is threatened whenever a **fault** occurs in any of the computers on which it runs or in the network that connects them.
- **Security model:** the openness of distributed systems exposes them to attack by both external and internal agents.

System Models

-
- 2.1 Introduction
 - 2.2 Architectural Models
 - 2.3 Fundamental Models**
 - 2.3.1 Interaction Model**



fourth edition

DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg



Interaction Model

- Distributed systems are composed of many processes interacting in complex ways.
- For example:
 - ▶ Multiple server processes may cooperate with one another to provide a service
 - ➔ Domain Name Service, which partitions and replicates its data at servers throughout the Internet
 - ▶ A set of peer processes may cooperate with one another to achieve a common goal
 - ➔ A voice conference system that distributes streams of audio data in a similar manner, but with strict real-time constraints.

Distributed Algorithm

- **Algorithm**: a sequence of steps to be taken in order to perform a desired computation.
- Distributed systems composed of multiple processes can be described by a **distributed algorithm**: a definition of the steps to be taken by each of the processes of which the system is composed, *including the transmission of messages between them*.
- **Messages** are transmitted between processes to transfer information between them and to coordinate their **activity**.

Some Assumptions

- The rate at which each process proceeds **cannot** in general be predicted.
- The timing of the transmission of messages **cannot** in general be predicted.
- Each process has its own state, consisting of the set of data that it can access and update.
- The state belonging to each process is completely private (that is, it cannot be accessed or updated by any other processes).

Factors Affecting Interacting Processes

- Communication performance is often a limiting characteristic.
- It is impossible to maintain a single global notion of time.

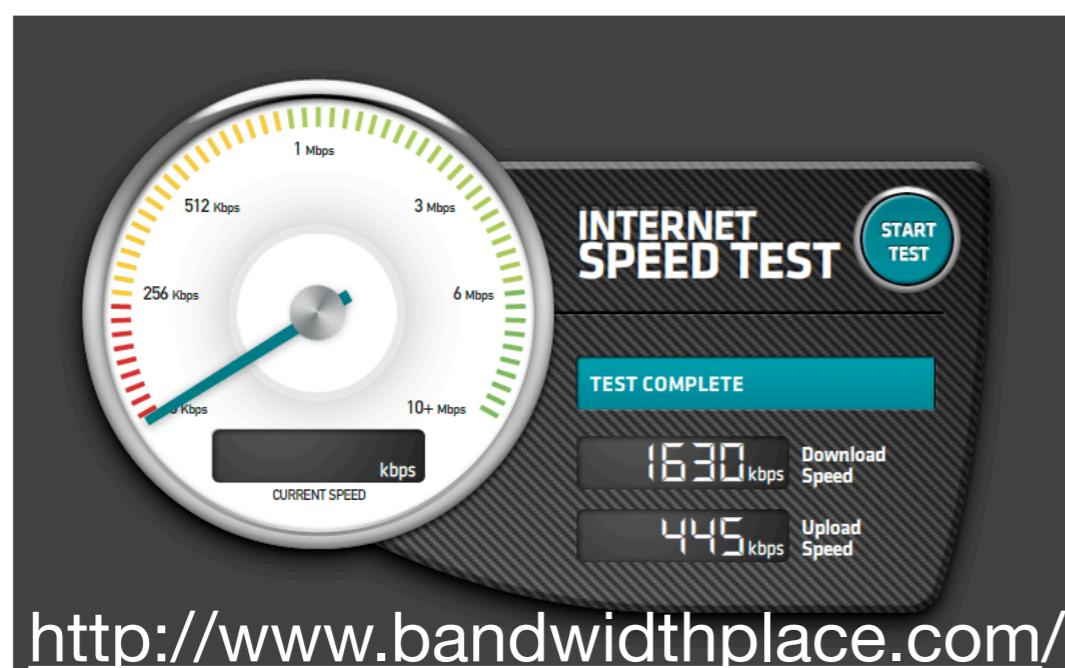


Performance of Communication Channels: Latency

- **Latency:** the delay between the start of a message's transmission from one process and the beginning of its receipt by another.
- The latency includes:
 - ▶ The **time taken for the first of a string of bits transmitted through the network to reach its destination.**
 - ▶ The **delay in accessing the network**, which increases significantly when the network is heavily loaded.
 - ▶ The **time taken by the operating system communication services** at both the sending and receiving processes, which varies according to the current load of the operating systems.

Performance of Communication Channels: Bandwidth

- The **bandwidth** of a computer network is the **total amount of information that can be transmitted over it in a given time**.
- Usually expressed in **bit/s** or multiples of it (kbit/s, Mbit/s, etc)
- When a large number of communication channels are using the same network, they have to share the available bandwidth.



56 kbit/s	Modem / Dialup
1.5 Mbit/s	ADSL Lite
1.544 Mbit/s	T1
10 Mbit/s	Ethernet
11 Mbit/s	Wireless 802.11b
44.736 Mbit/s	T3
54 Mbit/s	Wireless-G 802.11g
100 Mbit/s	Fast Ethernet
155 Mbit/s	OC3
300 Mbit/s	Wireless-N 802.11n
622 Mbit/s	OC12
1000 Mbit/s	Gigabit Ethernet
2.5 Gbit/s	OC48
9.6 Gbit/s	OC192
10 Gbit/s	10 Gigabit Ethernet

<http://www.bandwidthplace.com/>

Computer Clocks and Timing Events



- Each computer in a distributed system has its own internal clock, which can be used by local processes to obtain a value of the current time.
- Therefore, two processes running on different computers can associate timestamps with their events.
- However, even if two processes read their clocks at the same time, their local clocks may supply different time values.
- This is because computer clocks drift from perfect time and, more importantly, their drift rates differ from one another.
- **Clock drift rate:** relative amount that a computer clock differs from a perfect reference clock.

Two Variants of the Interaction Model

- In a distributed system it is hard to set time limits on the time taken for process execution, message delivery or clock drift.
- Two **opposite extreme positions** provide a pair of **simple models**:
 - ▶ **Synchronous distributed systems**: strong assumption of time
 - ▶ **Asynchronous distributed systems**: no assumptions about time

Synchronous Distributed System

- A distributed system in which the following **bounds** are defined:
 - ▶ the **time to execute each step** of a process has known lower and upper bounds
 - ▶ each **message transmitted** over a channel is **received** within a known bounded time
 - ▶ each process has a **local clock** whose **drift rate from real time** has a known bound

Asynchronous Distributed System

- A distributed system in which there are no bounds on:
 - ▶ process execution speeds: each step may take an arbitrarily long time
 - ▶ message transmission delays: a message may be received after an arbitrarily long time
 - ▶ clock drift rates: the drift rate of a clock is arbitrary
- This exactly models the Internet, in which there is no intrinsic bound on server or network load and therefore on how long it takes, for example, to transfer a file using ftp, or to receive an email message.
- *Any solution that is valid for an asynchronous distributed system is also valid for a synchronous one. Why? What about the contrary?*

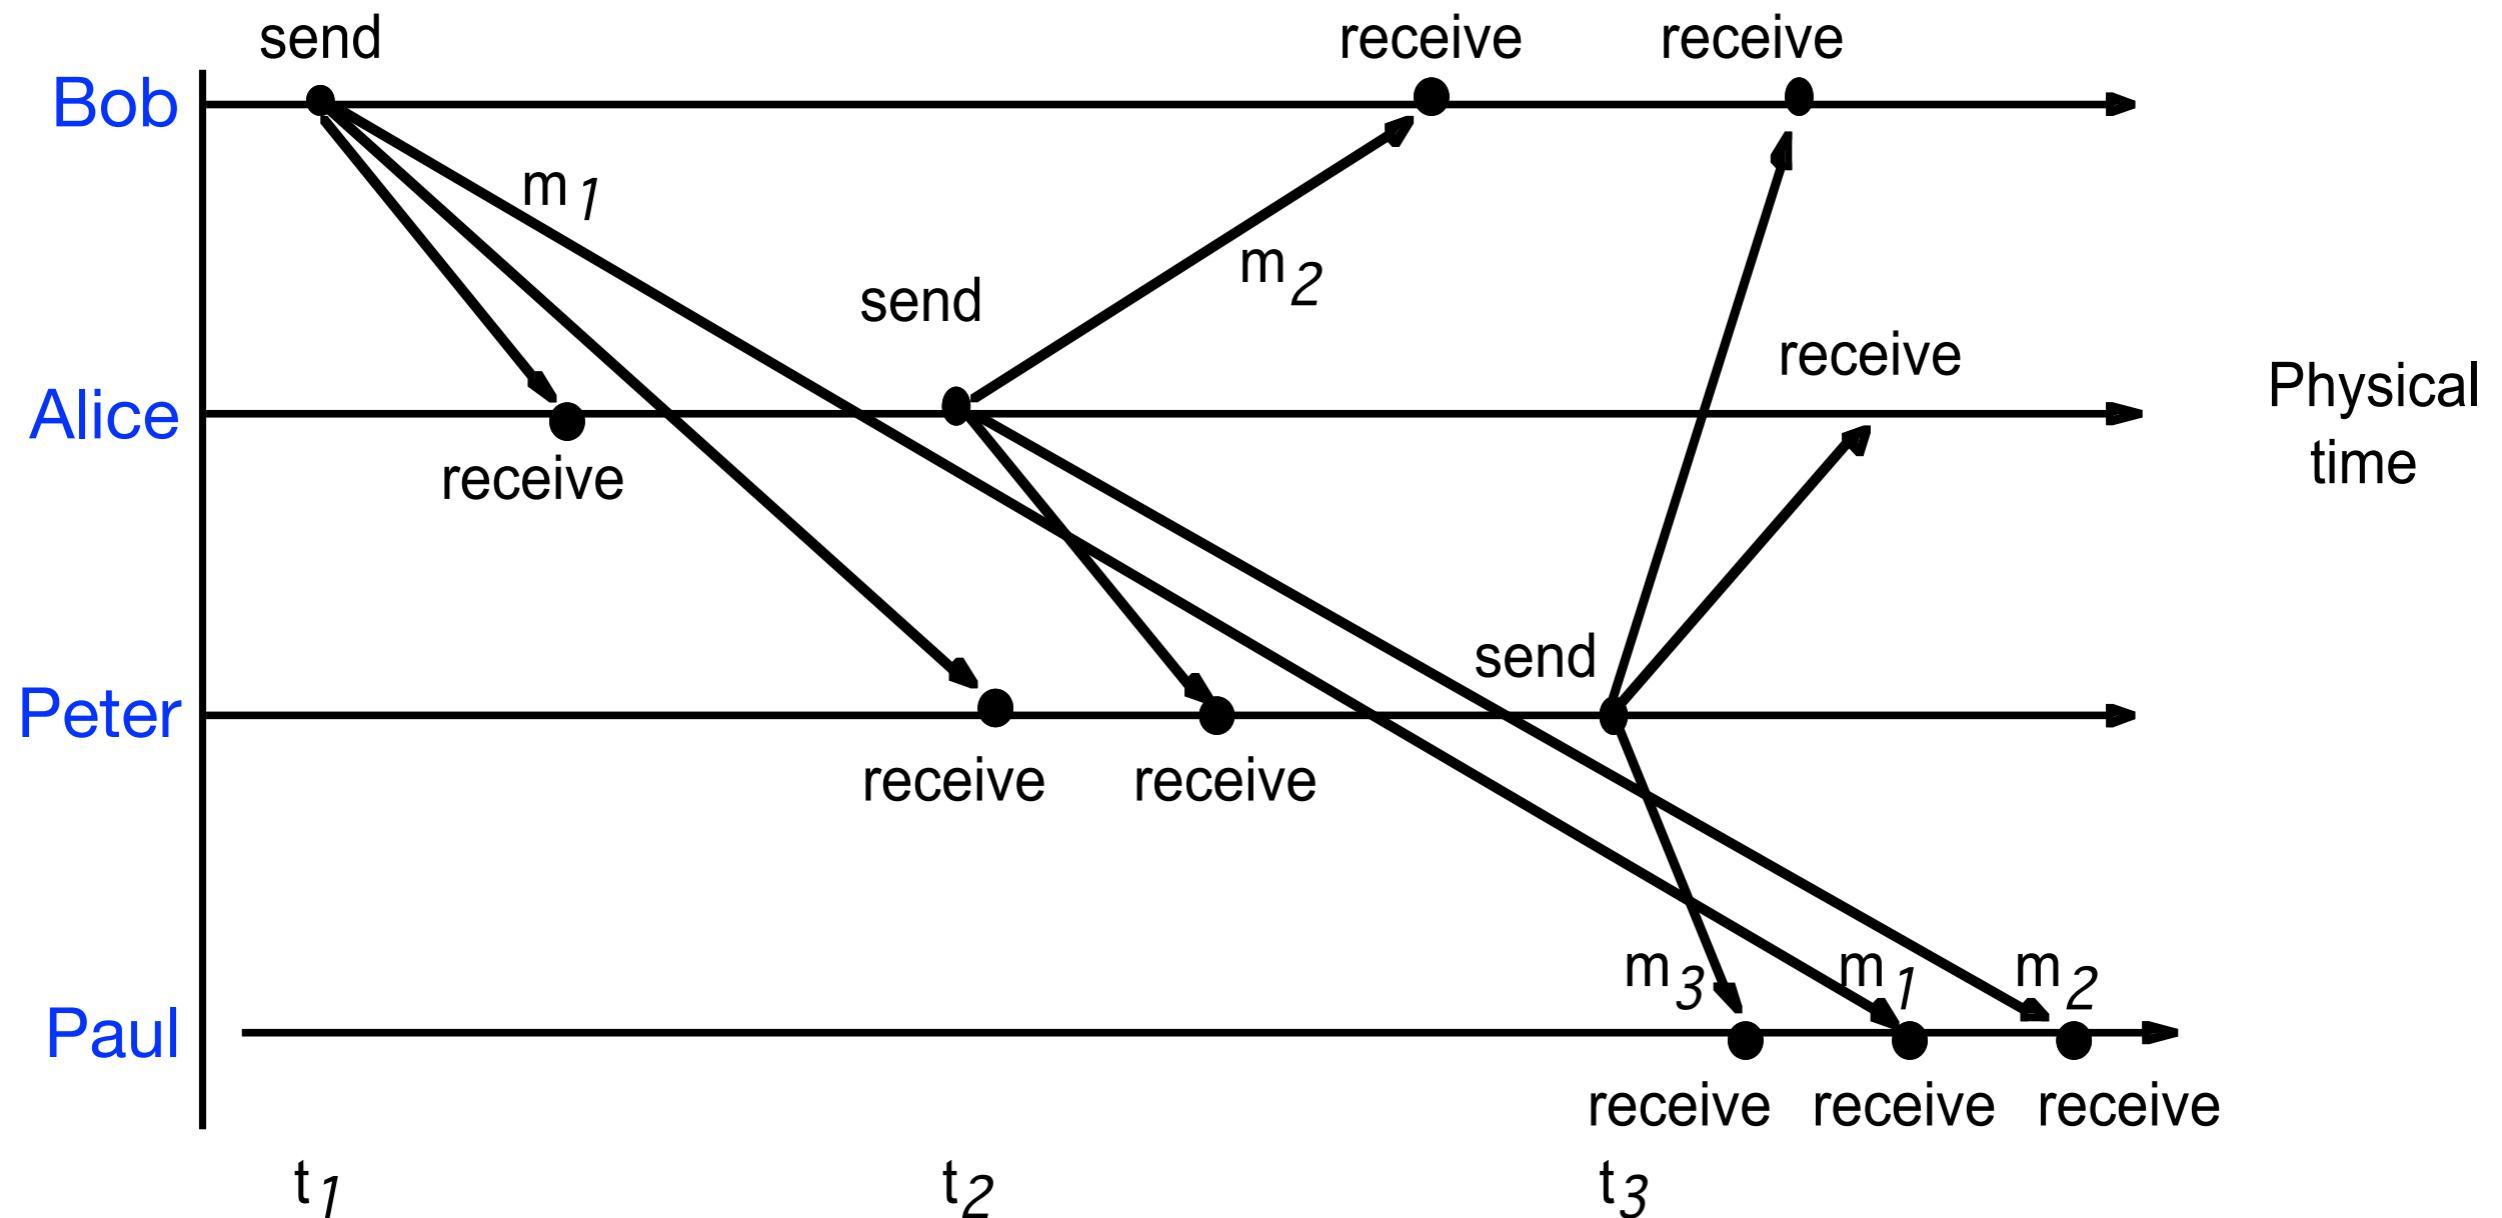
Event Ordering

- In many cases, we are interested in knowing whether an event (sending or receiving a message) at one process occurred *before*, *after* or *concurrently* with another event at another process.
- The execution of a system can be described in terms of events and their ordering despite the lack of accurate clocks.

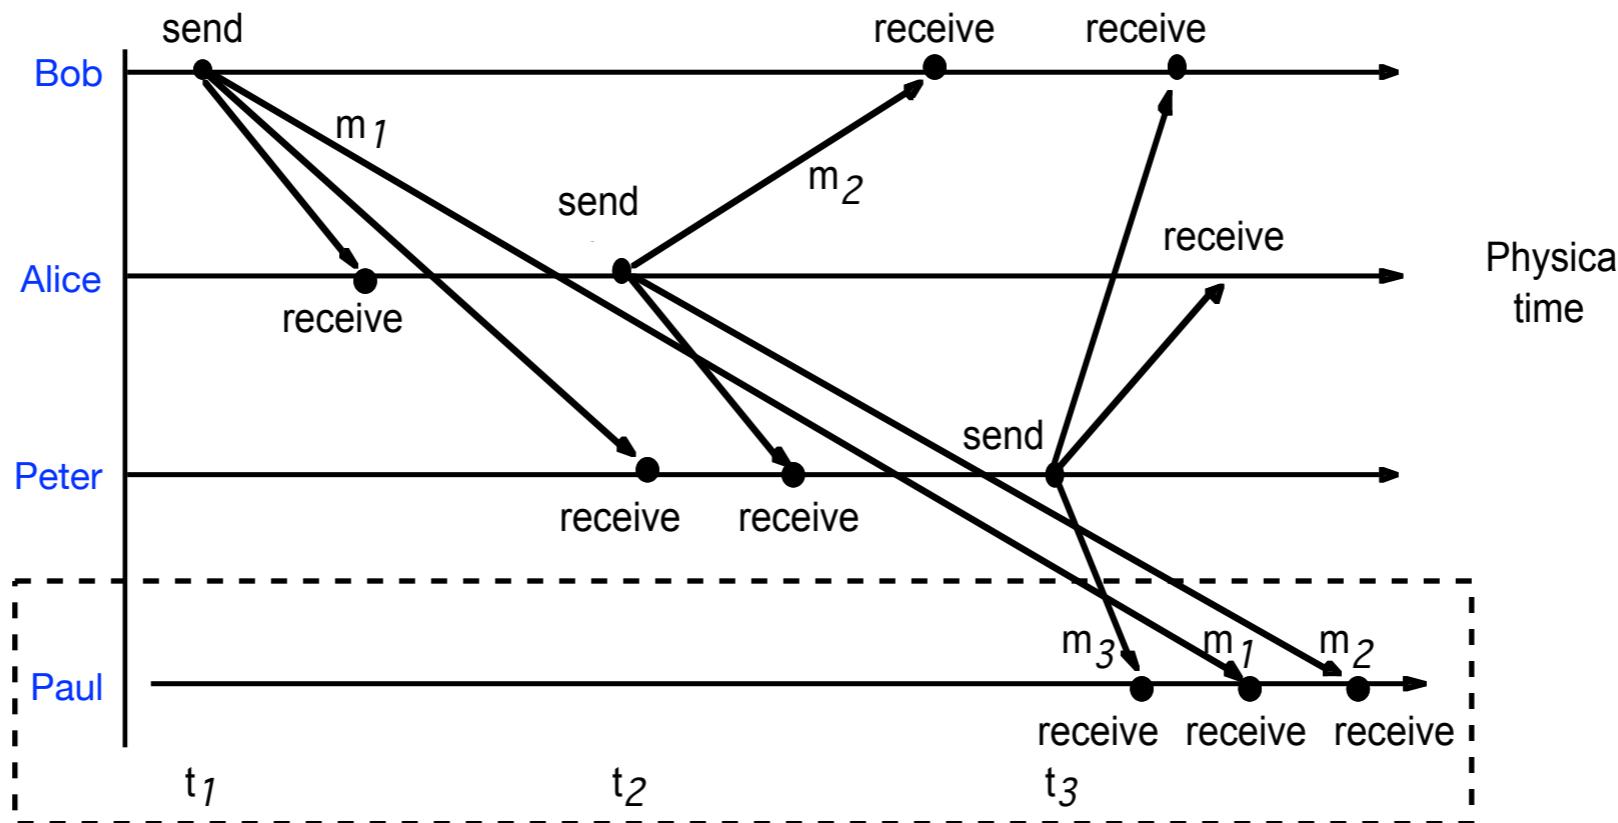
Event Ordering

- In many cases, we are interested in knowing whether an event (sending or receiving a message) at one process occurred *before*, *after* or *concurrently* with another event at another process.
- The execution of a system can be described in terms of events and their ordering despite the lack of accurate clocks.
- Example [Real-Time Ordering of Events]: consider the following set of exchanges between a group of email users Bob, Alice, Peter, and Paul on a mailing list:
 1. Bob sends a message with the subject *Meeting*
 2. Alice and Peter reply by sending a message with the subject *Re: Meeting*

Example: Real-Time Ordering of Events



Example: Real-Time Ordering of Events (cont.)



Inbox	
From	Subject
Peter	Re:meeting
Bob	Meeting
Alice	Re: Meeting

System Models

-
- 2.1 Introduction
 - 2.2 Architectural Models
 - 2.3 Fundamental Models
 - 2.3.2 Failure Model



fourth edition

DISTRIBUTED SYSTEMS CONCEPTS AND DESIGN

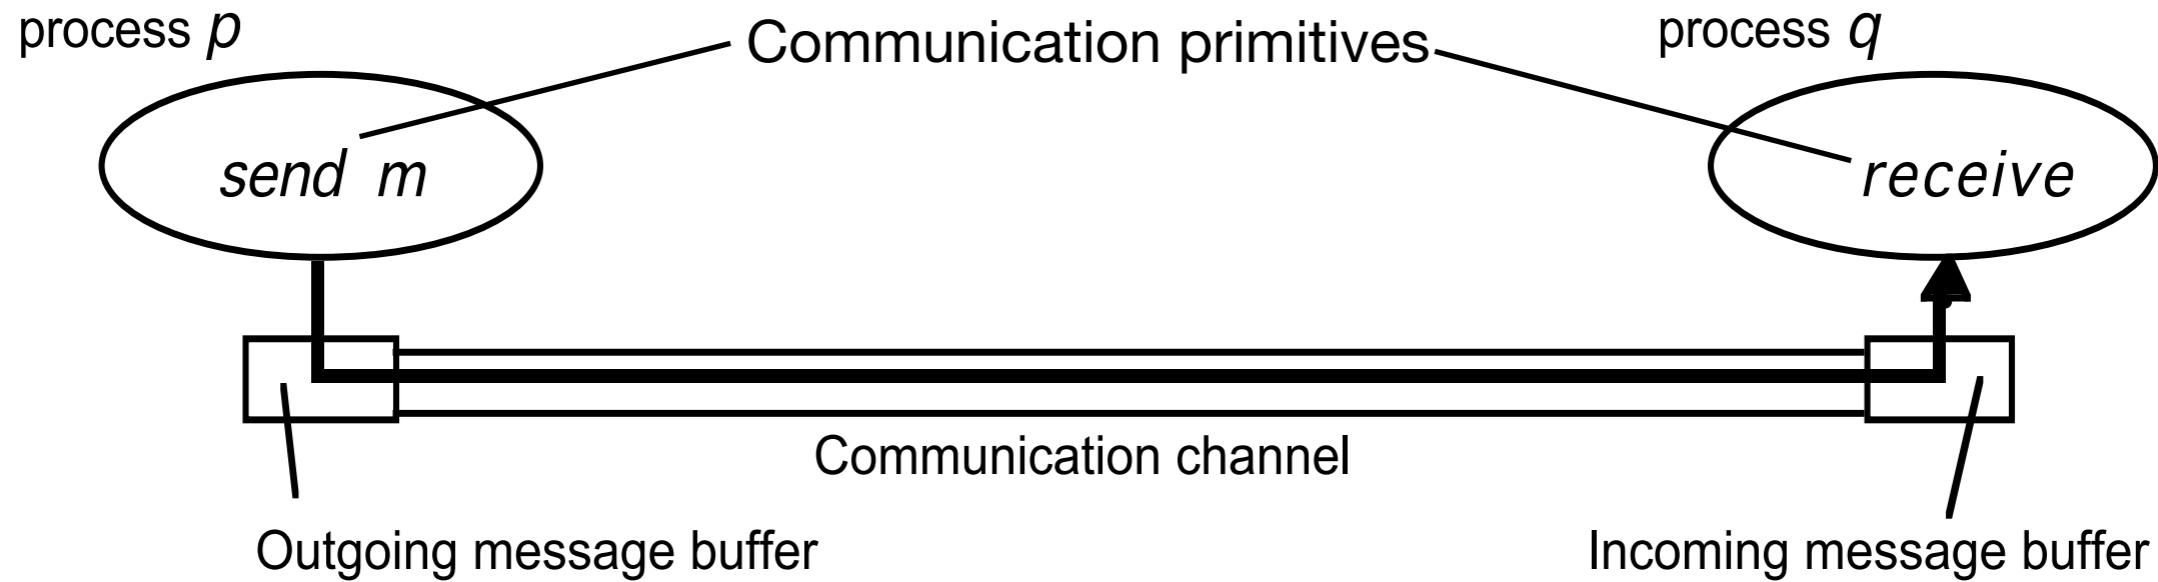
George Coulouris
Jean Dollimore
Tim Kindberg



Failure Model

- In a distributed system **both processes and communication channels may fail** (that is, they may depart from what is considered to be correct or desirable behavior).
- The **failure model** defines the ways in which failures may occur in order to provide an understanding of the effects of failures.
- Example of **taxonomy of failures** [Hadzilacos and Toueg, 1994]:
 - ▶ **Omission failures**: a process or communication channel fails to perform actions that it is supposed to do
 - ▶ **Arbitrary failures**: any type of error may occur
 - ▶ **Timing failures**: applicable in synchronous distributed systems

[Failure Model] Omission Failures



Class of failure	Affects	Description
Crash	Process	Process halts prematurely and remain halted.
Omission	Channel	A msg inserted in an outgoing msg buffer never arrives at the other end's incoming message buffer.
Send-omission	Process	A process completes a send, but the message is not put in its outgoing message buffer.
Receive-omission	Process	A message is put in a process's incoming message buffer, but that process does not receive it.

[Failure Model] Arbitrary Failures

- The term *arbitrary* or *Byzantine* failure is used to describe the *worst possible failure semantics*, in which *any type of error may occur*.
- *Arbitrary failure of a process*: the process arbitrarily omits intended processing steps or takes unintended processing steps.
- *Communication channel arbitrary failures*: message contents may be corrupted or non-existent messages may be delivered or real messages may be delivered more than once.

Class of failure	Affects	Description
Arbitrary (Byzantine)	Process or channel	Process/channel exhibits arbitrary behaviour: it may send/transmit arbitrary messages at arbitrary times, commit omissions; a process may stop or take an incorrect step.

[Failure Model] Timing Failures

- Timing failures are applicable in **synchronous** distributed systems, where time limits are set on process execution time, message delivery time and clock drift rate.

Class of failure	Affects	Description
Clock	Process	Process's local clock exceeds the bounds on its rate of drift from real time.
Performance	Process	Process exceeds the bounds on the interval between two steps.
Performance	Channel	A message's transmission takes longer than the stated bound.

[Failure Model] Timing Failures

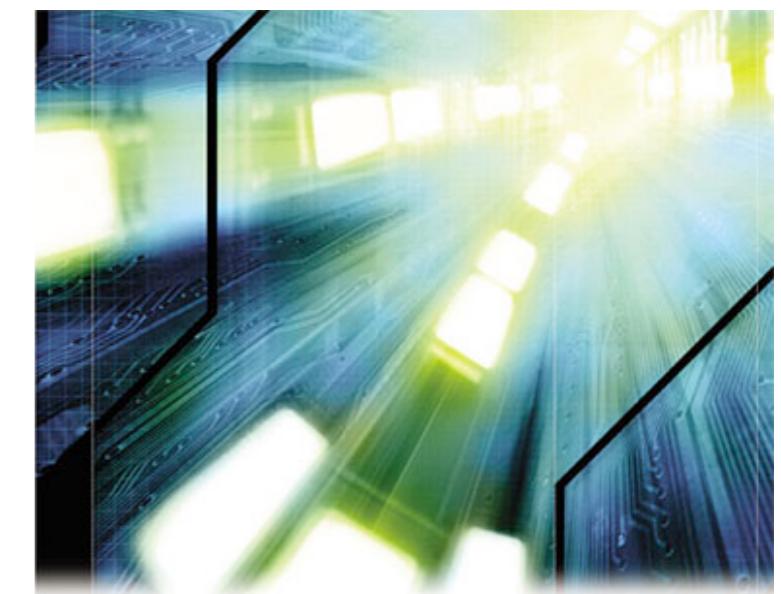
- Timing failures are applicable in **synchronous** distributed systems, where time limits are set on process execution time, message delivery time and clock drift rate.

Class of failure	Affects	Description
Clock	Process	Process's local clock exceeds the bounds on its rate of drift from real time.
Performance	Process	Process exceeds the bounds on the interval between two steps.
Performance	Channel	A message's transmission takes longer than the stated bound.

- In an **asynchronous** distributed systems, an overloaded server may respond too slowly, but *we cannot say that it has a timing failure since no guarantee has been offered.*

System Models

-
- 2.1 Introduction
 - 2.2 Architectural Models
 - 2.3 Fundamental Models**
 - 2.3.3 Security Model**



DISTRIBUTED SYSTEMS
CONCEPTS AND DESIGN

George Coulouris
Jean Dollimore
Tim Kindberg

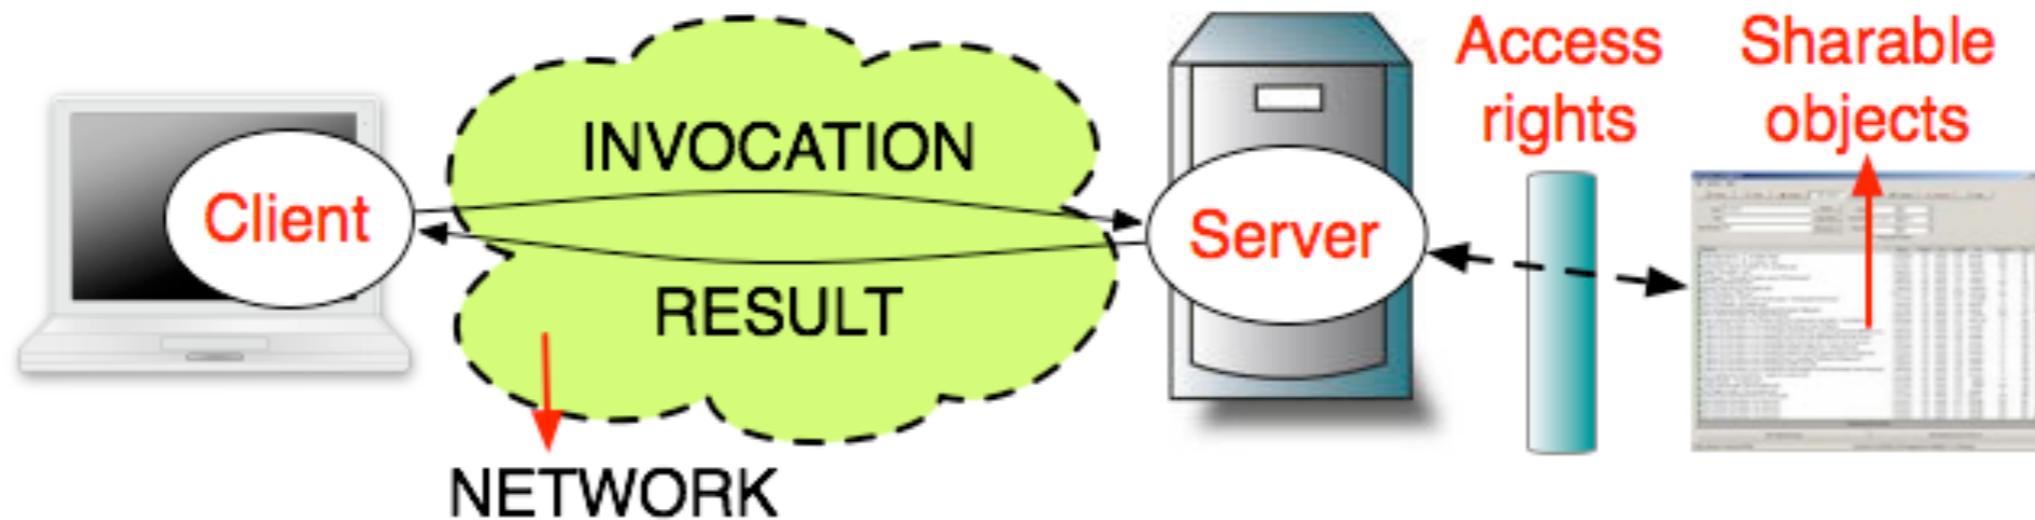


Security Model

- Motivating factor for distributed systems: sharing of resources.
- Architectural model: distributed system described in terms of processes encapsulating (sharable) objects and providing access to them through interactions with other processes.
- This model provides the basis for the security model:

The security of a distributed system can be achieved by securing the processes and the channels used for their interactions and by protecting the objects that they encapsulate against unauthorized access.

[Security Model] Protecting Objects



- Objects are intended to be used in different ways by different users:
 - ▶ some objects may hold a user's private data (such as the mailbox)
 - ▶ other objects may hold shared data such as web pages.
- **Access rights** specify **who is allowed to perform the operations on an object** (for instance, who is allowed to read and write its state).

[Security Model] Security Threats

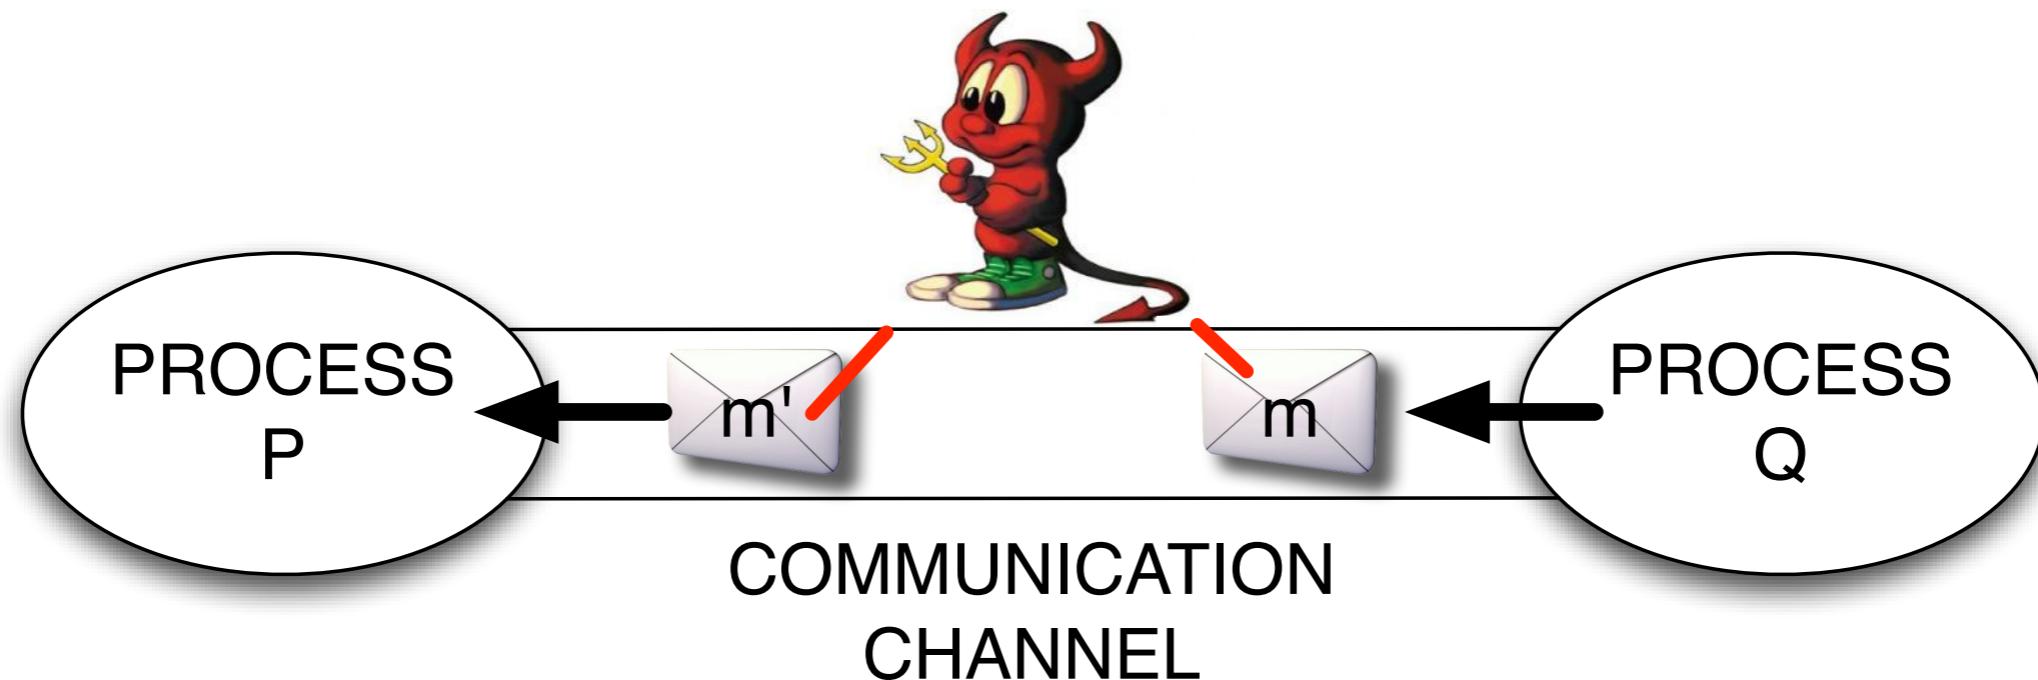
- A **security threat** is a potential violation of security.
- To model security threats, we postulate an **enemy** (or **adversary**) that is *capable of sending any message to any process and reading or copying any message between a pair of processes*.



- The threats from potential enemy are discussed under the headings **threats to processes**, **threats to communication channels** and **denial of service**.

[Security Model] Threats to Processes

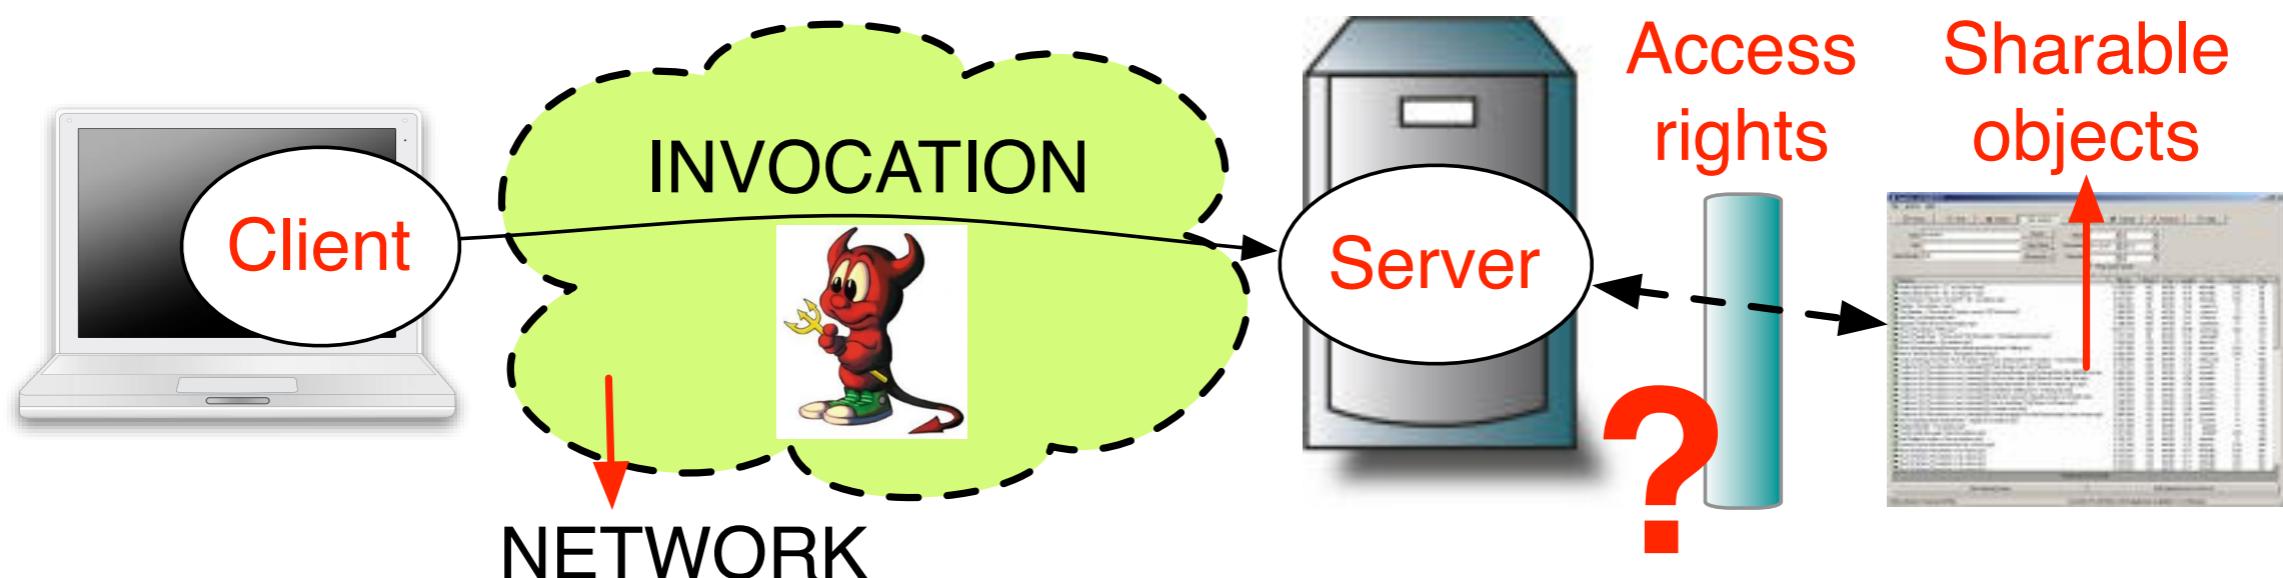
- A process receives a message from another process in the distributed system, and **it is not able to determine the identity of the sender**.



m' = message with a forged source address

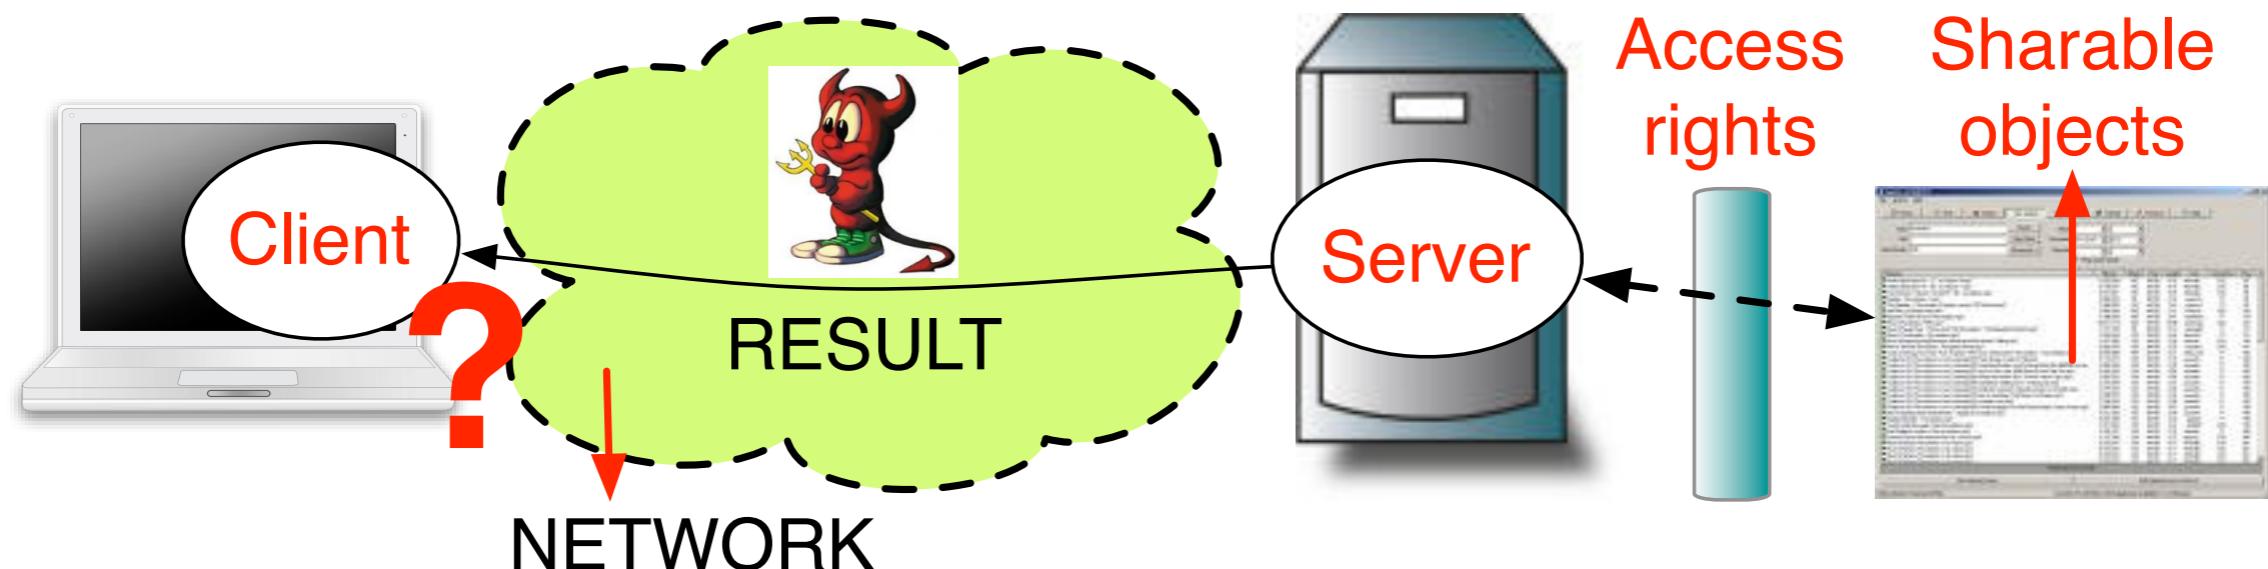
[Security Model] Threats to Processes

- This lack of reliable knowledge is a threat to the correct functioning of both servers and clients:
 - ▶ Server: without reliable knowledge of the sender's identity, a server cannot tell whether to perform the operation or reject it.



[Security Model] Threats to Processes

- This lack of reliable knowledge is a threat to the correct functioning of both servers and clients:
 - ▶ Client: when a client receives the result of an invocation from a server, it cannot necessarily tell whether the source of the result message is from the intended server or from an enemy, perhaps “spoofing” the mail server.



[Security Model] Threats to Communication Channels

- An **enemy can copy, alter or inject messages** as they travel across the network.
- Such attacks present **a threat to the privacy and integrity of information** as it travels over the network and **to the integrity of the system**.
 - ▶ Example: a result message containing a user's mail item might be revealed to another user or it might be altered to say something quite different.

[Security Model] Threats to Communication Channels

- An **enemy can copy, alter or inject messages** as they travel across the network.
- Such attacks present **a threat to the privacy and integrity of information** as it travels over the network and **to the integrity of the system**.
 - ▶ Example: a result message containing a user's mail item might be revealed to another user or it might be altered to say something quite different.
- Another form of **attack** is the attempt to **save copies of messages and to reply them at a later time**, making it possible to reuse the same message over and over again.
 - ▶ Example: someone could benefit by resending an invocation message requesting a transfer of a sum of money from one bank account to another.

[Security Model] Denial of Service

- This is a form of attack in which the **enemy interferes with the activities of authorized users** by making **excessive and pointless invocations** on services or message transmissions in a network, resulting in **overloading of physical resources** (network bandwidth, server processing capacity, ...).
- Intentions: **delaying or preventing actions by other users.**



[Security Model] Denial of Service

- This is a form of attack in which the **enemy interferes with the activities of authorized users** by making **excessive and pointless invocations** on services or message transmissions in a network, resulting in **overloading of physical resources** (network bandwidth, server processing capacity, ...).
- Intentions: **delaying or preventing actions by other users.**



THU AUG
6TH

Example: on August 6, 2009, **Twitter** was shut down for hours due to a DoS attack:

Ongoing denial-of-service attack 21 hours ago

We are defending against a denial-of-service attack, and will update status again shortly.

Update: the site is back up, but we are continuing to defend against and recover from this attack.

Update (9:46a): As we recover, users will experience some longer load times and slowness. This includes timeouts to API clients. We're working to get back to 100% as quickly as we can.

Update (4:14p): Site latency has continued to improve, however some web requests continue to fail. This means that some people may be unable to post or follow from the website.