

UNIT - III

75

PEER TO PEER SERVICES AND FILE SYSTEM

Peer-to-peer systems represent a paradigm for the construction of distributed systems and applications in which data and computational resources are contributed by many hosts on the Internet, all of which participate in the provision of uniform service.

Peer-to-peer middleware systems are emerging that have the capacity to share computing resources, storage and data present in computers at the edge of the Internet on a global scale.

Napster and its Legacy:

Napster architecture included centralized indexes but users supplied the files which were stored and accessed on their personal computers. Napster was shut down as a result of legal proceeding that instituted against the operators of the Napster service by the owner of the copyright in some of the material that was made available on it.

Peer-to-peer middleware:

A key problem in the design of Peer-to-peer applications is to provide a mechanism to enable clients to access data resources quickly.

and dependably wherever they are located
throughout the network.

76

Peer-to-peer middleware systems are designed specifically to meet the need for automatic placement and subsequent location of the distributed objects managed by peer-to-peer systems and applications.

Functional requirements

The function of peer-to-peer middleware is to simplify the construction of services that are implemented across many hosts in a widely distributed network.

Other important requirements include the ability to add new resources and remove them at will and to add hosts to the service and remove them.

Non-Functional Requirements

To perform effectively, peer-to-peer middleware must also address the following non-functional requirement.

* Global scalability

* Load Balancing

76

77

* Optimization for local interactions between neighbouring peers

* Accommodating to highly dynamic host availability

* Security of data in an environment with heterogeneous trust

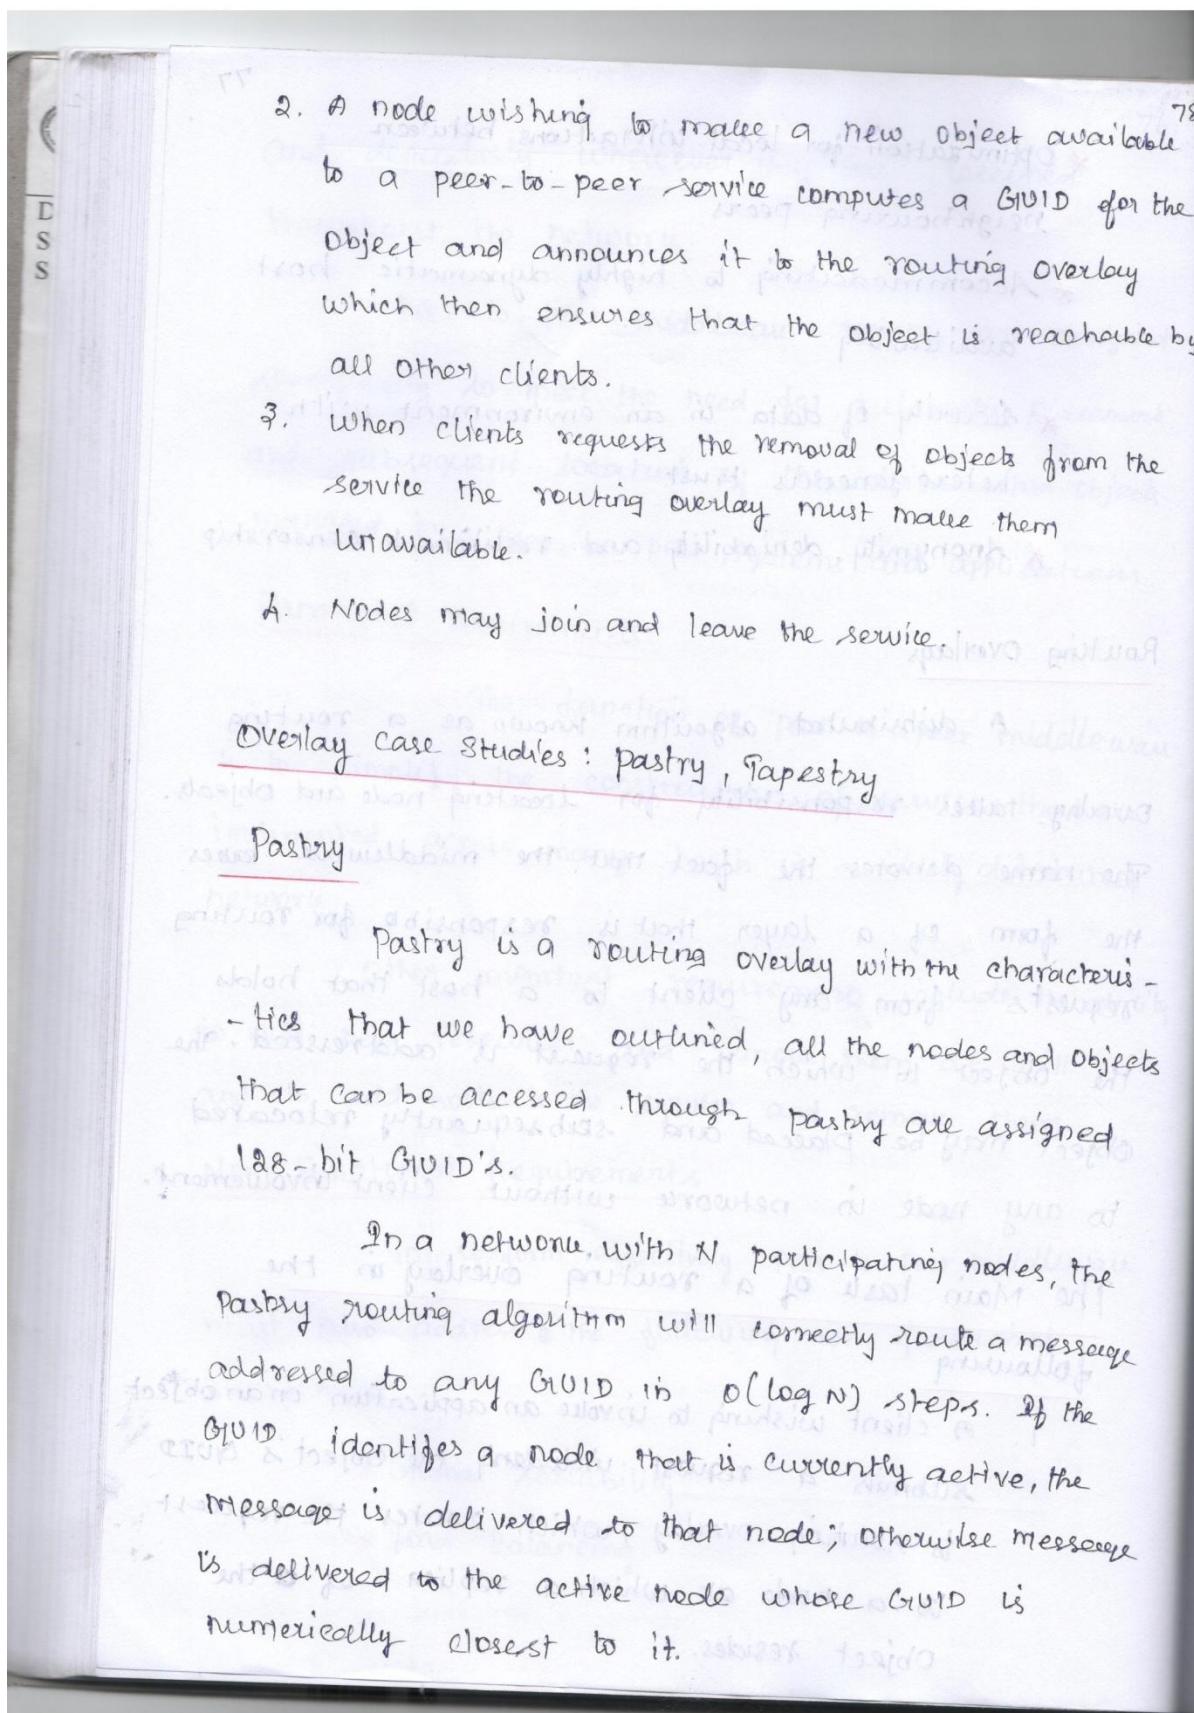
* Anonymity, deniability and resilience to censorship

Routing Overlays

A distributed algorithm known as a routing overlay takes responsibility for locating node and objects. The name denotes the fact that the middleware takes the form of a layer that is responsible for routing requests from any client to a host that holds the object to which the request is addressed. The object may be placed and subsequently relocated to any node in network without client involvement.

The Main task of a routing overlay is the following

1. A client wishing to invoke an application on an object submits a request including the object's GUID to routing overlay, which routes the request to a node at which a replica of the object resides.



78

available
for the
delay
table by
from the

79

Routing steps involve the use of an underlying transport protocol to transfer the message to a Pastry node that is closer to its destination.

Pastry uses a locality metric based on network distance in the underlying network to select appropriate neighbours when setting up the routing table used at each node.

Routing algorithm:

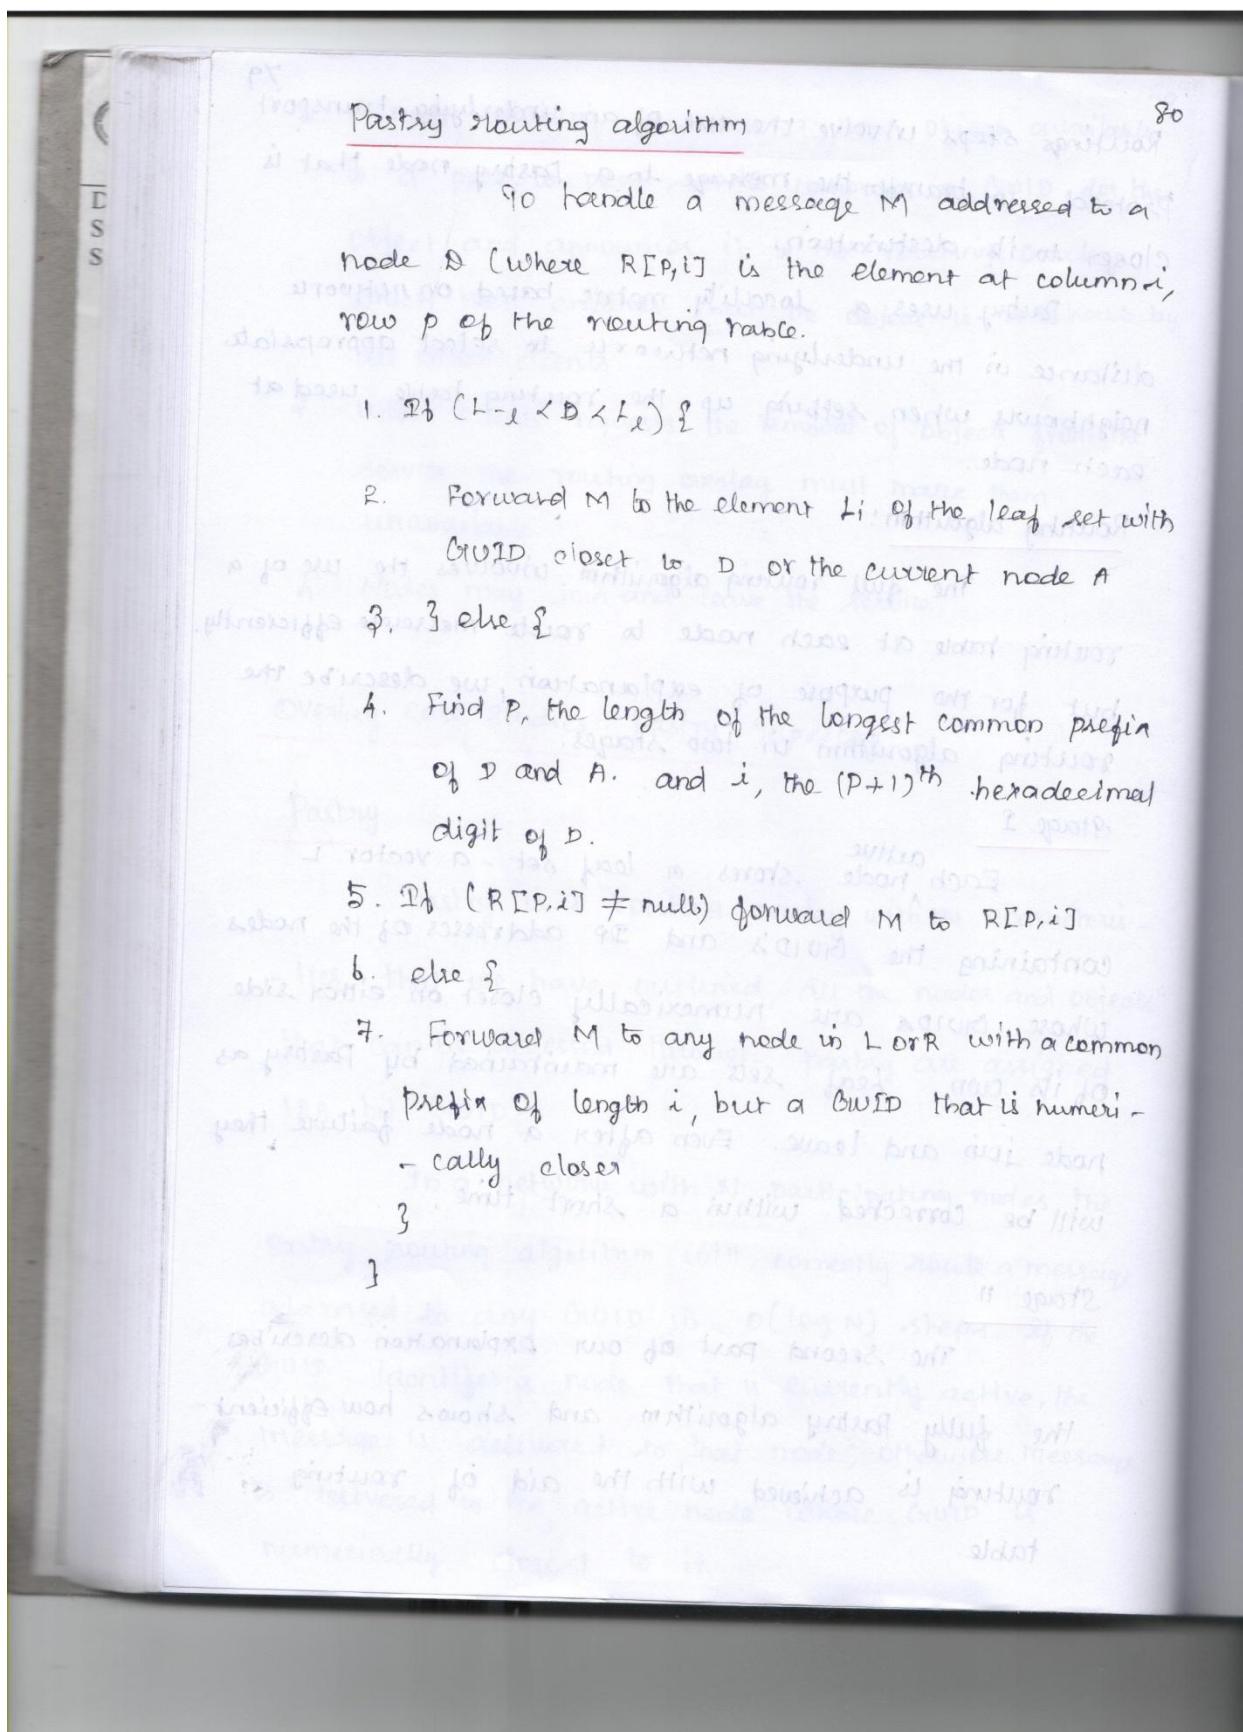
The full routing algorithm involves the use of a routing table at each node to route messages efficiently. but for the purpose of explanation, we describe the routing algorithm in two stages.

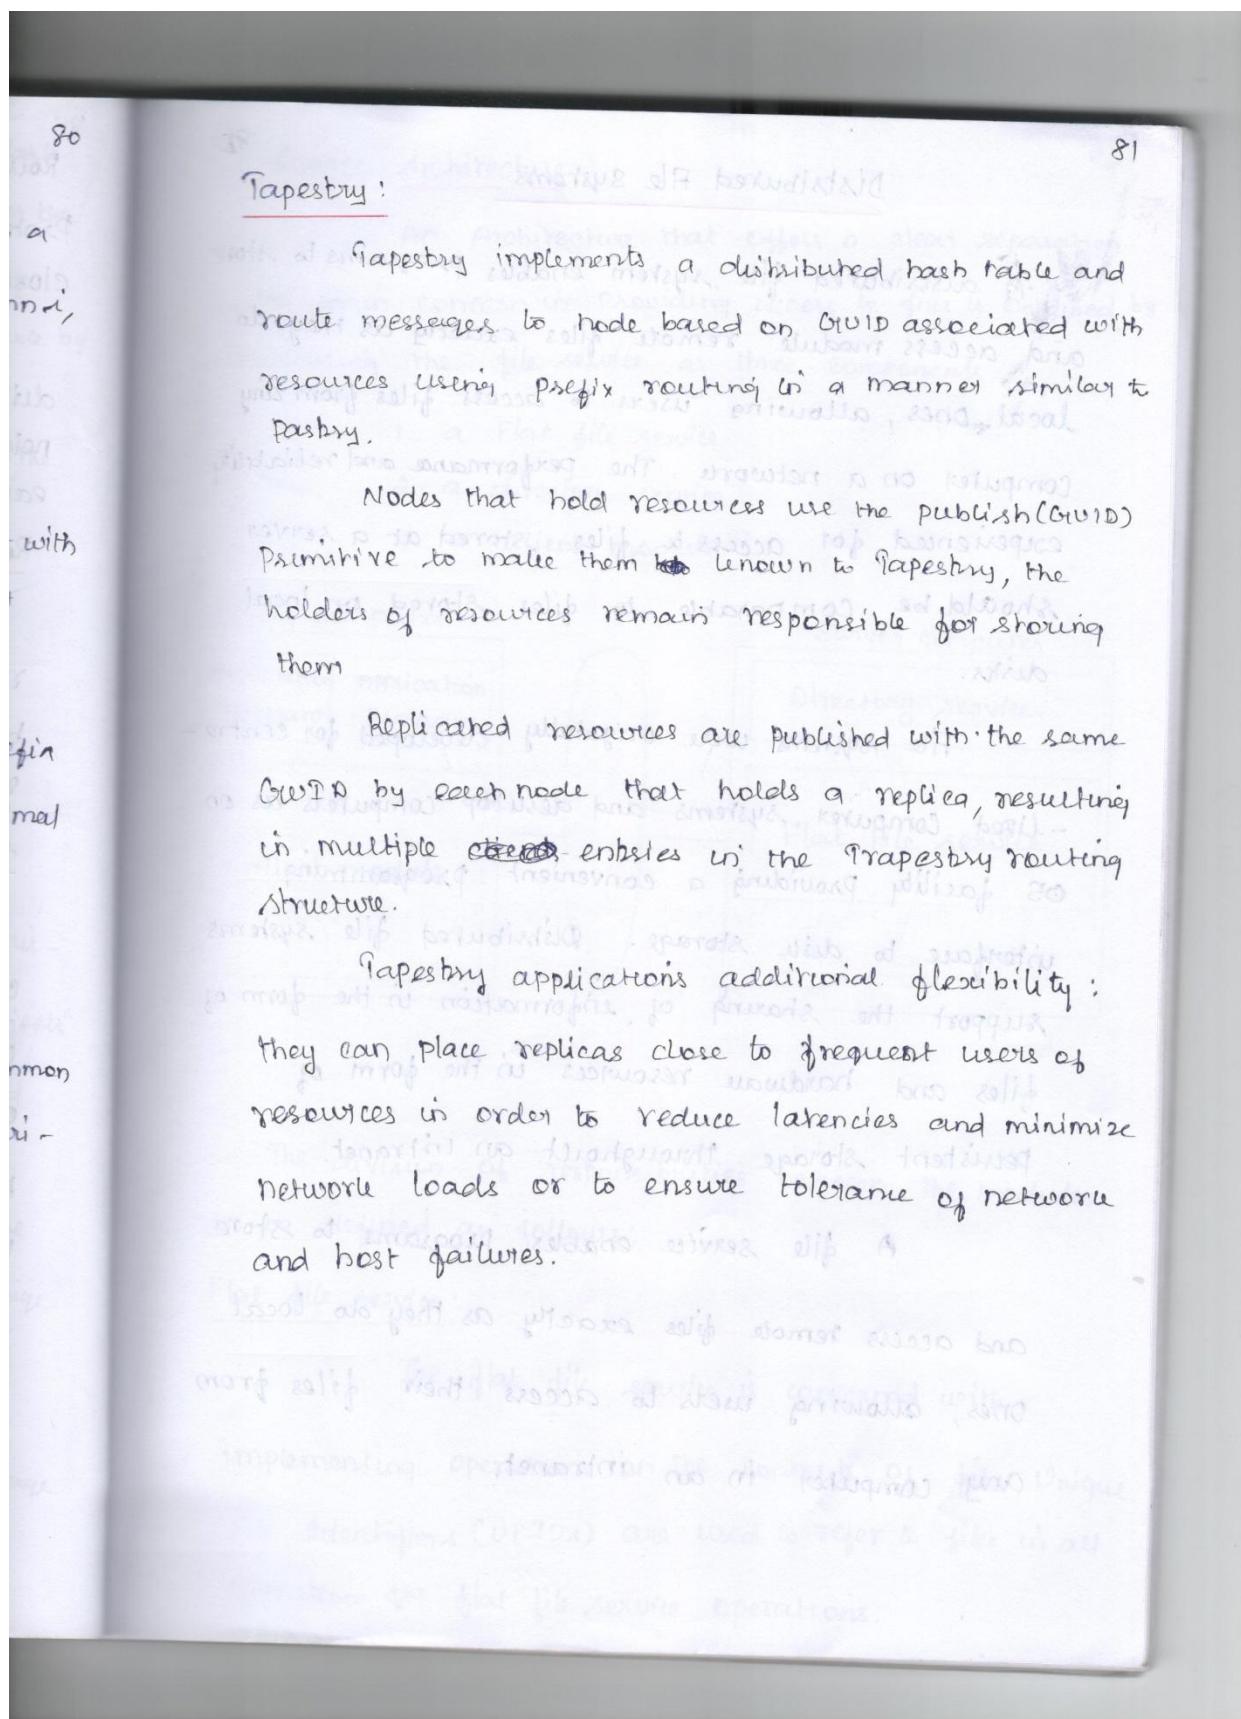
Stage I

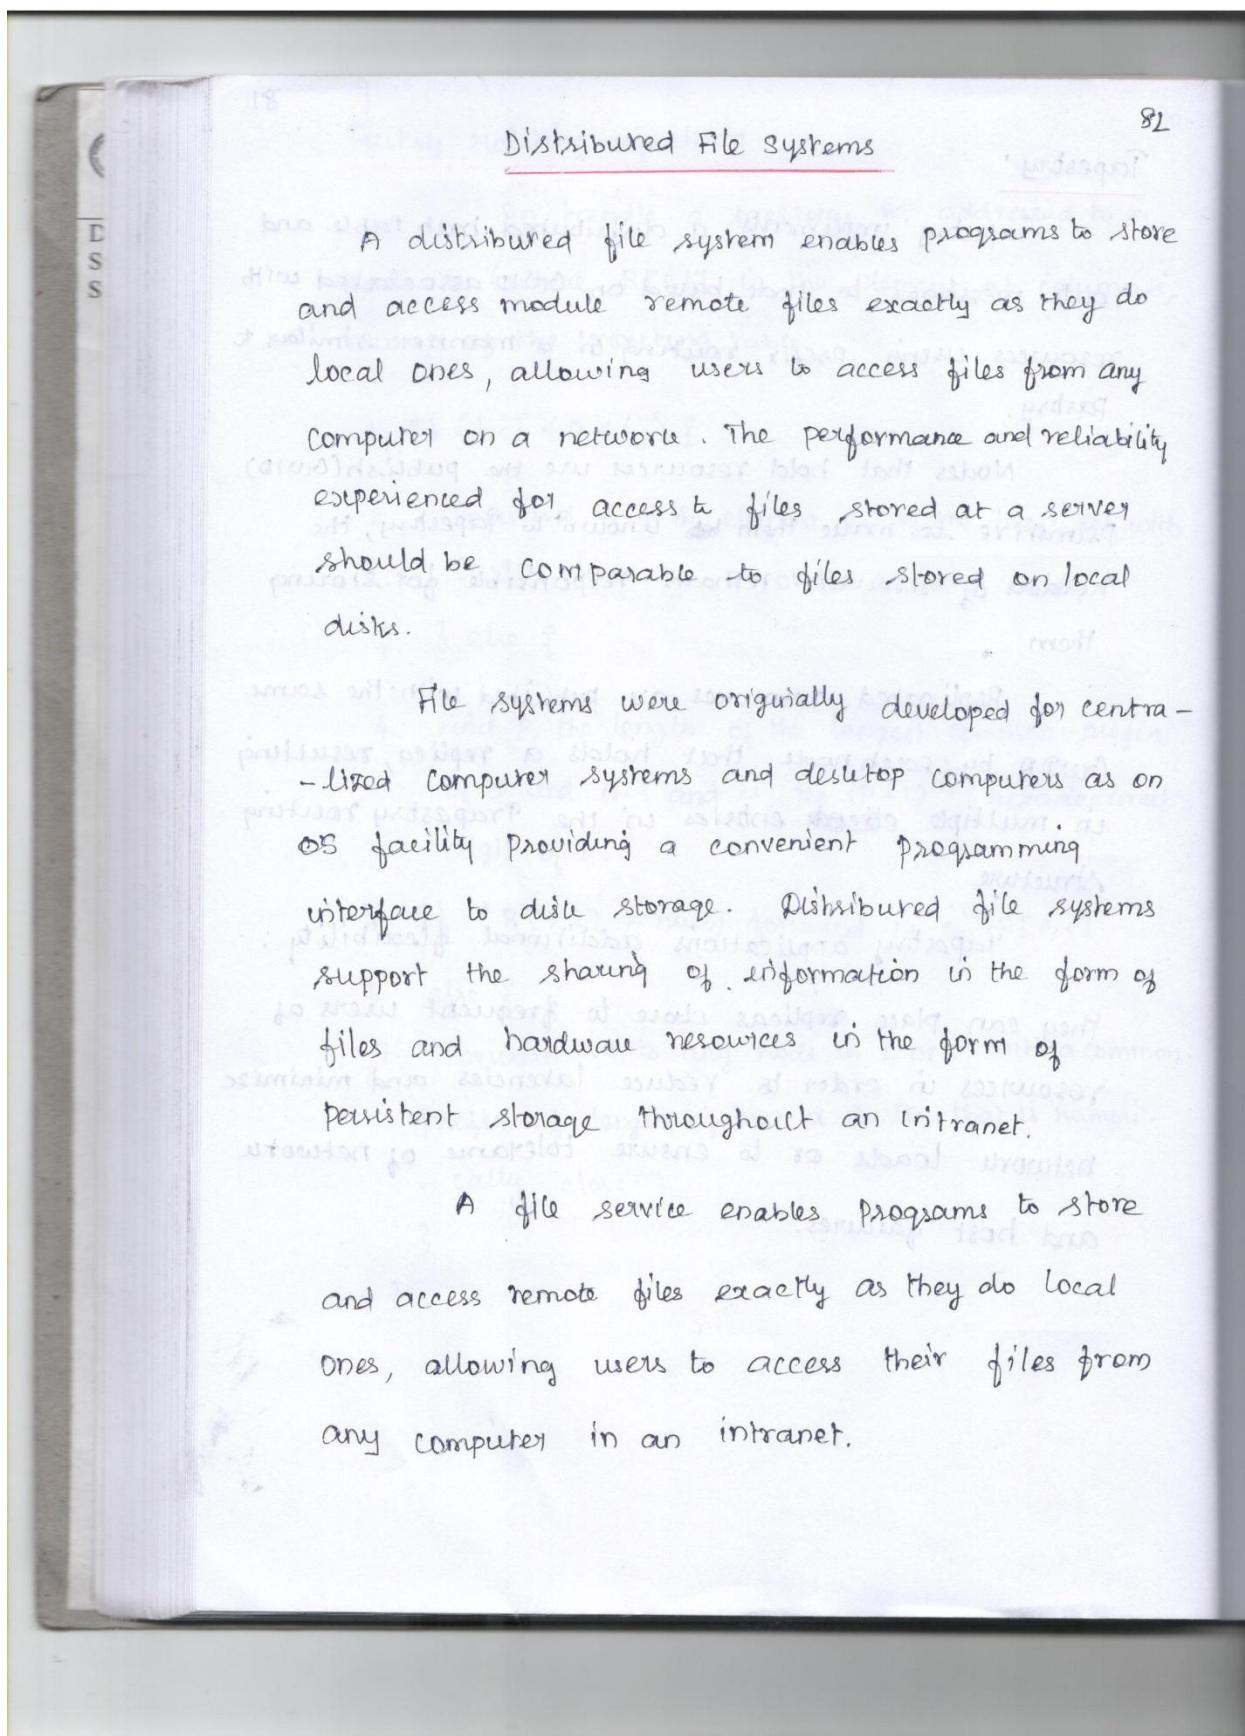
Each node stores an active leaf set - a vector L containing the GVID's and IP addresses of the nodes whose GVIDs are numerically closest on either side of its own. Leaf sets are maintained by Pastry as node join and leave. Even after a node failure they will be corrected within a short time.

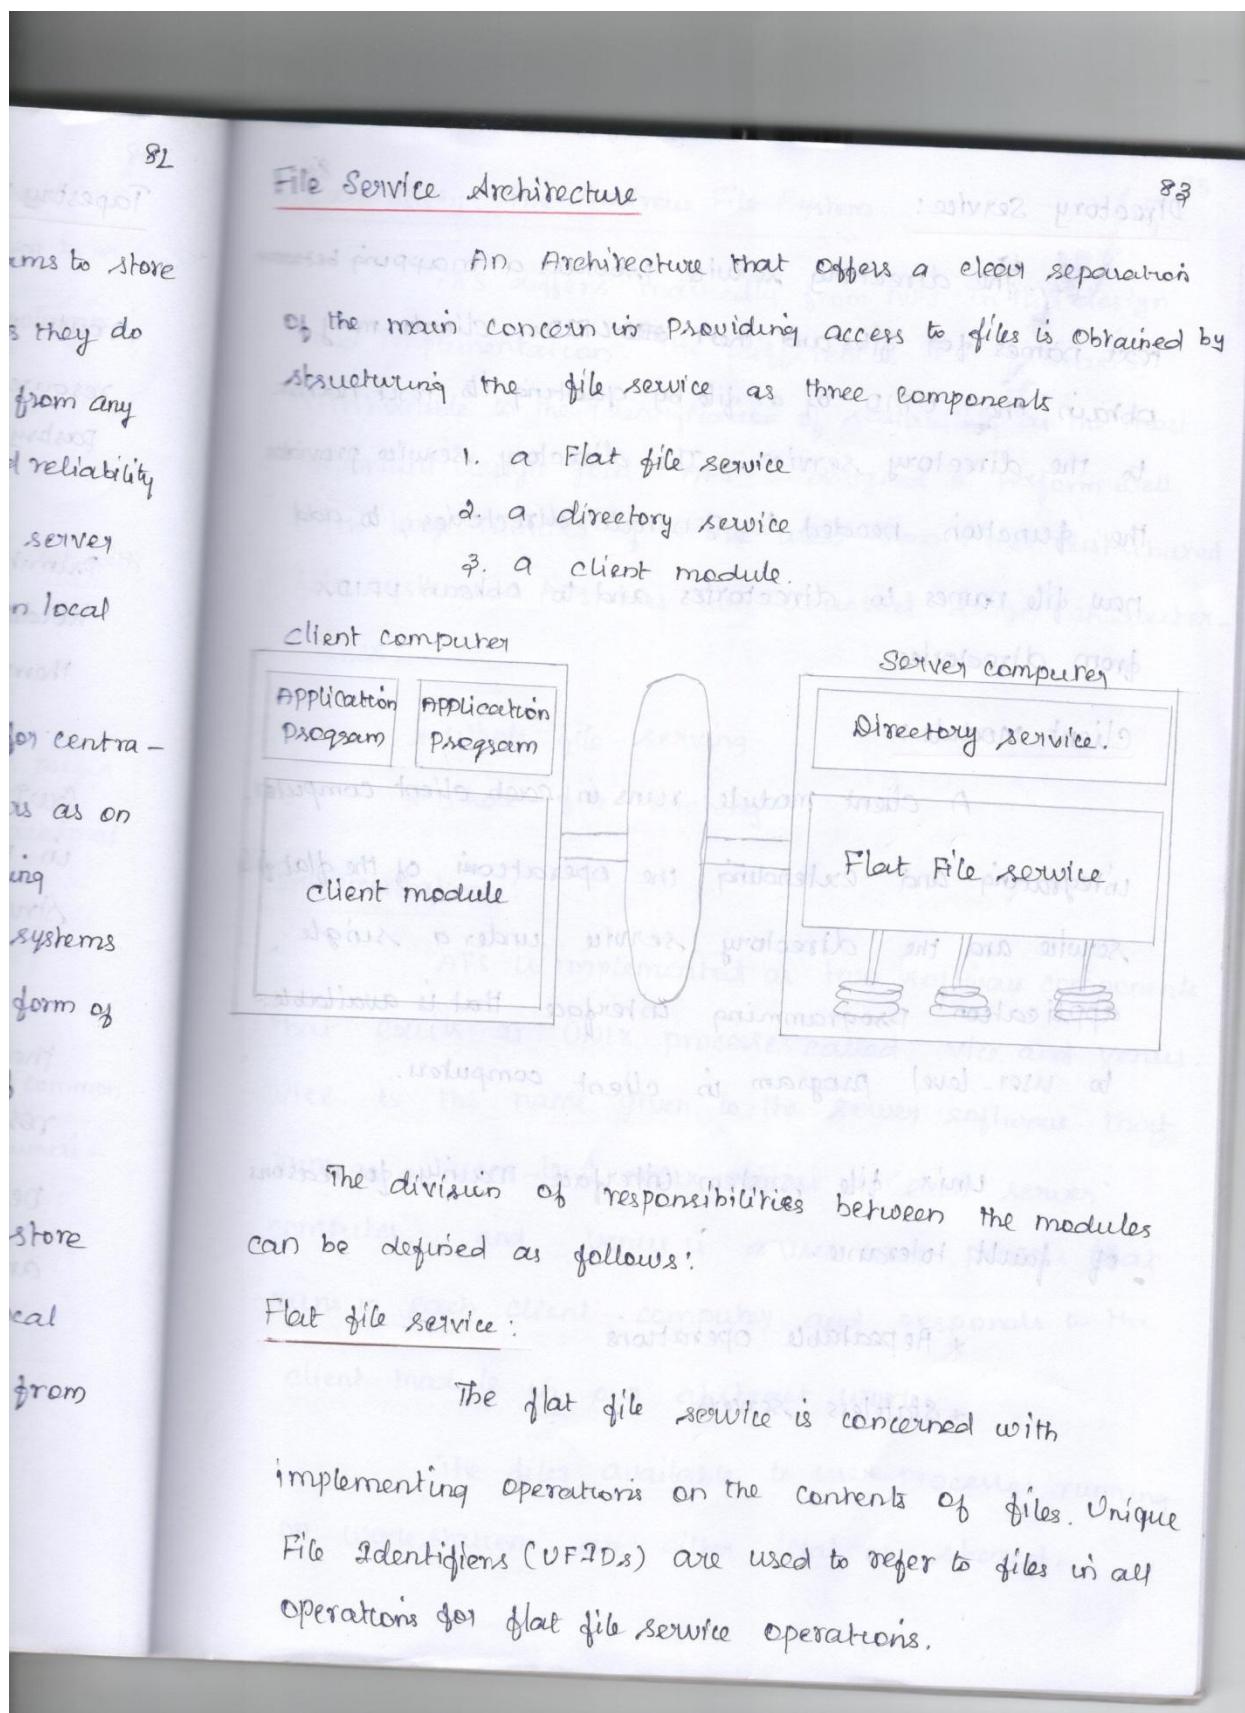
Stage II

The second part of our explanation describes the fully Pastry algorithm and shows how efficient routing is achieved with the aid of routing table.









The division of responsibilities between the modules can be defined as follows:

Flat file service:

The flat file service is concerned with implementing operations on the contents of files. Unique File Identifiers (UFIDs) are used to refer to files in all operations for flat file service operations.

Directory Service:

The directory service provides a mapping between text names for files and their ~~UFIDs~~ UIDs. Clients may obtain the UID of a file by quoting its text name to the directory service. The directory service provides the function needed to generate directories, to add new file names to directories and to obtain UIDs from directories.

Client Module:

A client module runs in each client computer, integrating and extending the operations of the flat file service and the directory service under a single application programming interface that is available to user-level programs in client computers.

Unix file system interface mainly for reasons of fault tolerance.

- * Repeatable operations
- * Stateless servers.

Case Study: The Andrew File Systems

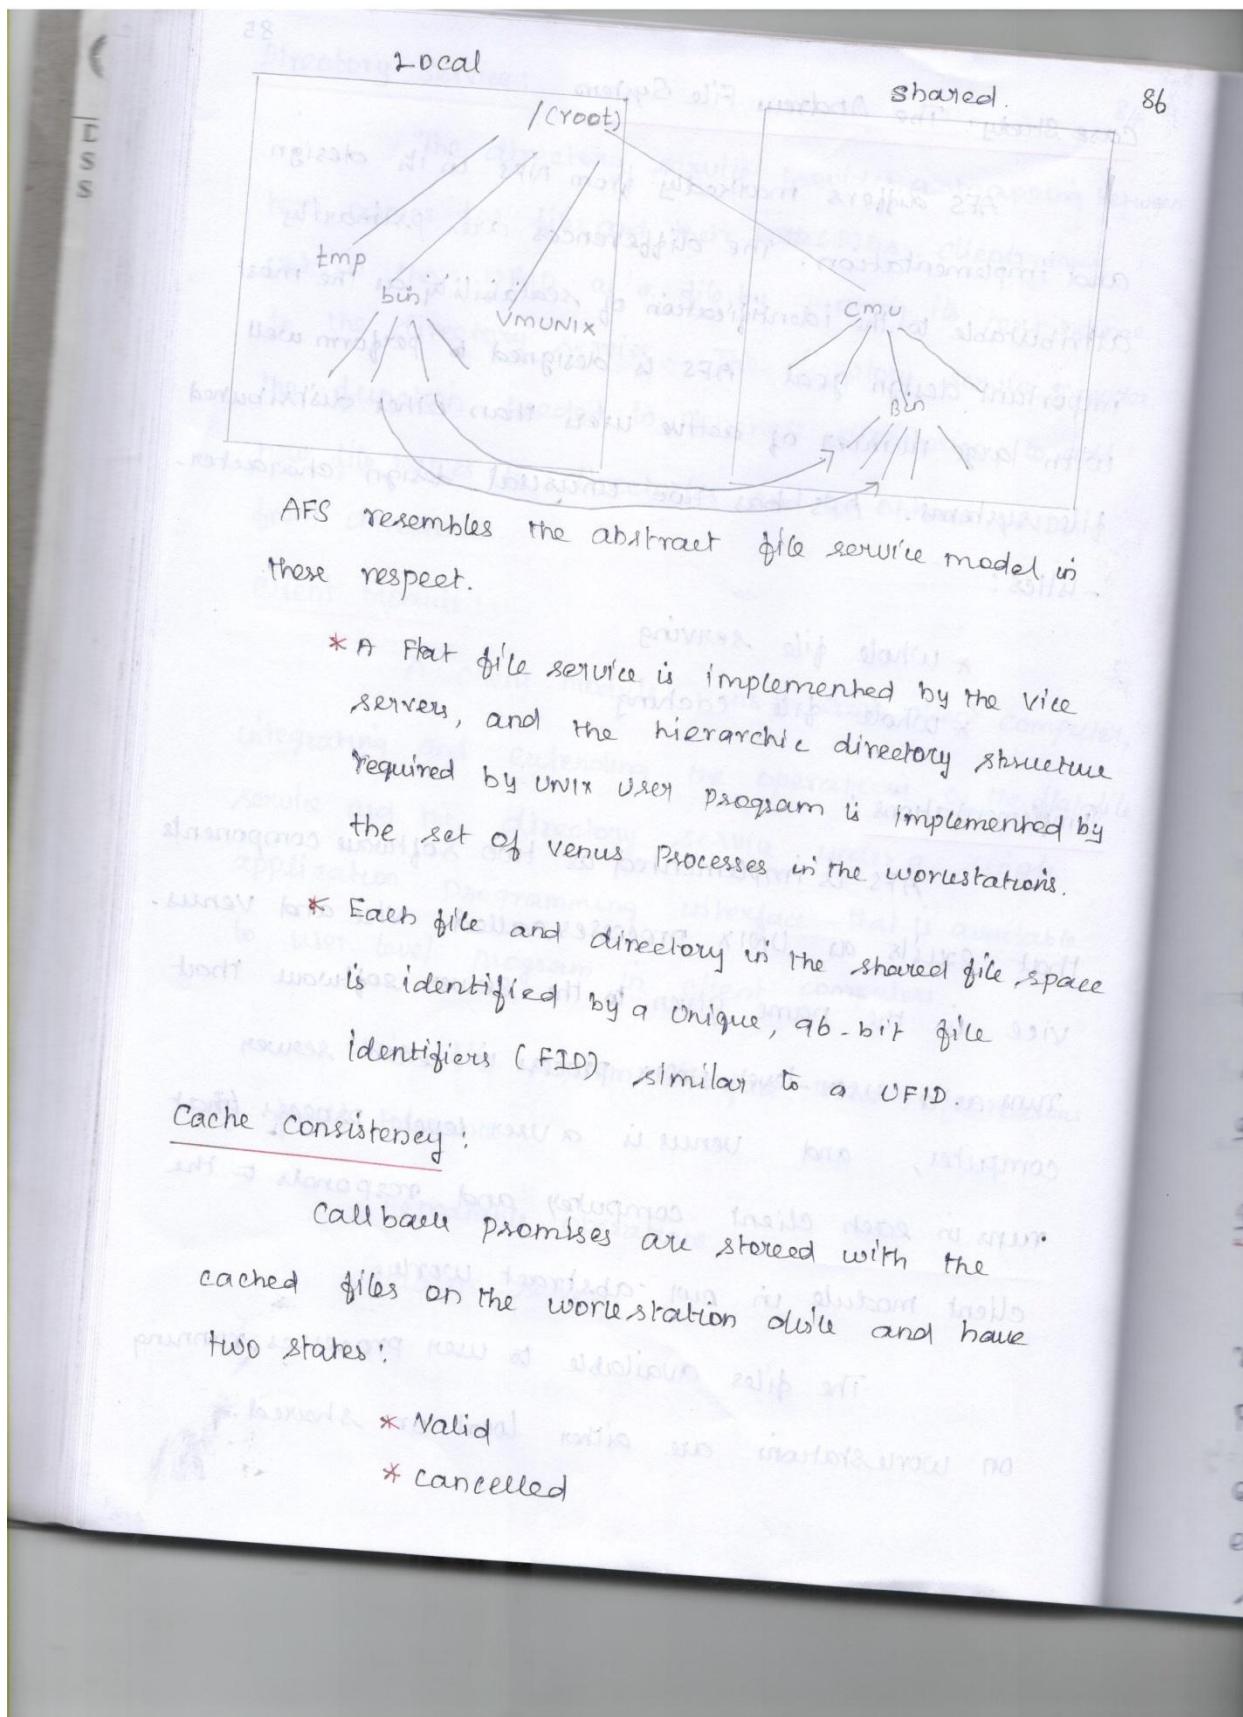
AFS differs markedly from NFS in its design and implementation. The differences are primarily attributable to the identification of scalability as the most important design goal. AFS is designed to perform well with large number of active users than other distributed file systems. AFS has two unusual design characteristics:

- * Whole file serving
- * Whole file caching

Implementations:

AFS is implemented as two software components that exists as UNIX processes called Vice and Venus. Vice is the name given to the server software that runs as a user-level UNIX process in each server computer, and Venus is a user level process that runs in each client computer and responds to the client module in our abstract work.

The files available to user processes running on workstations are either local or shared.



86 88 87

update Semantics

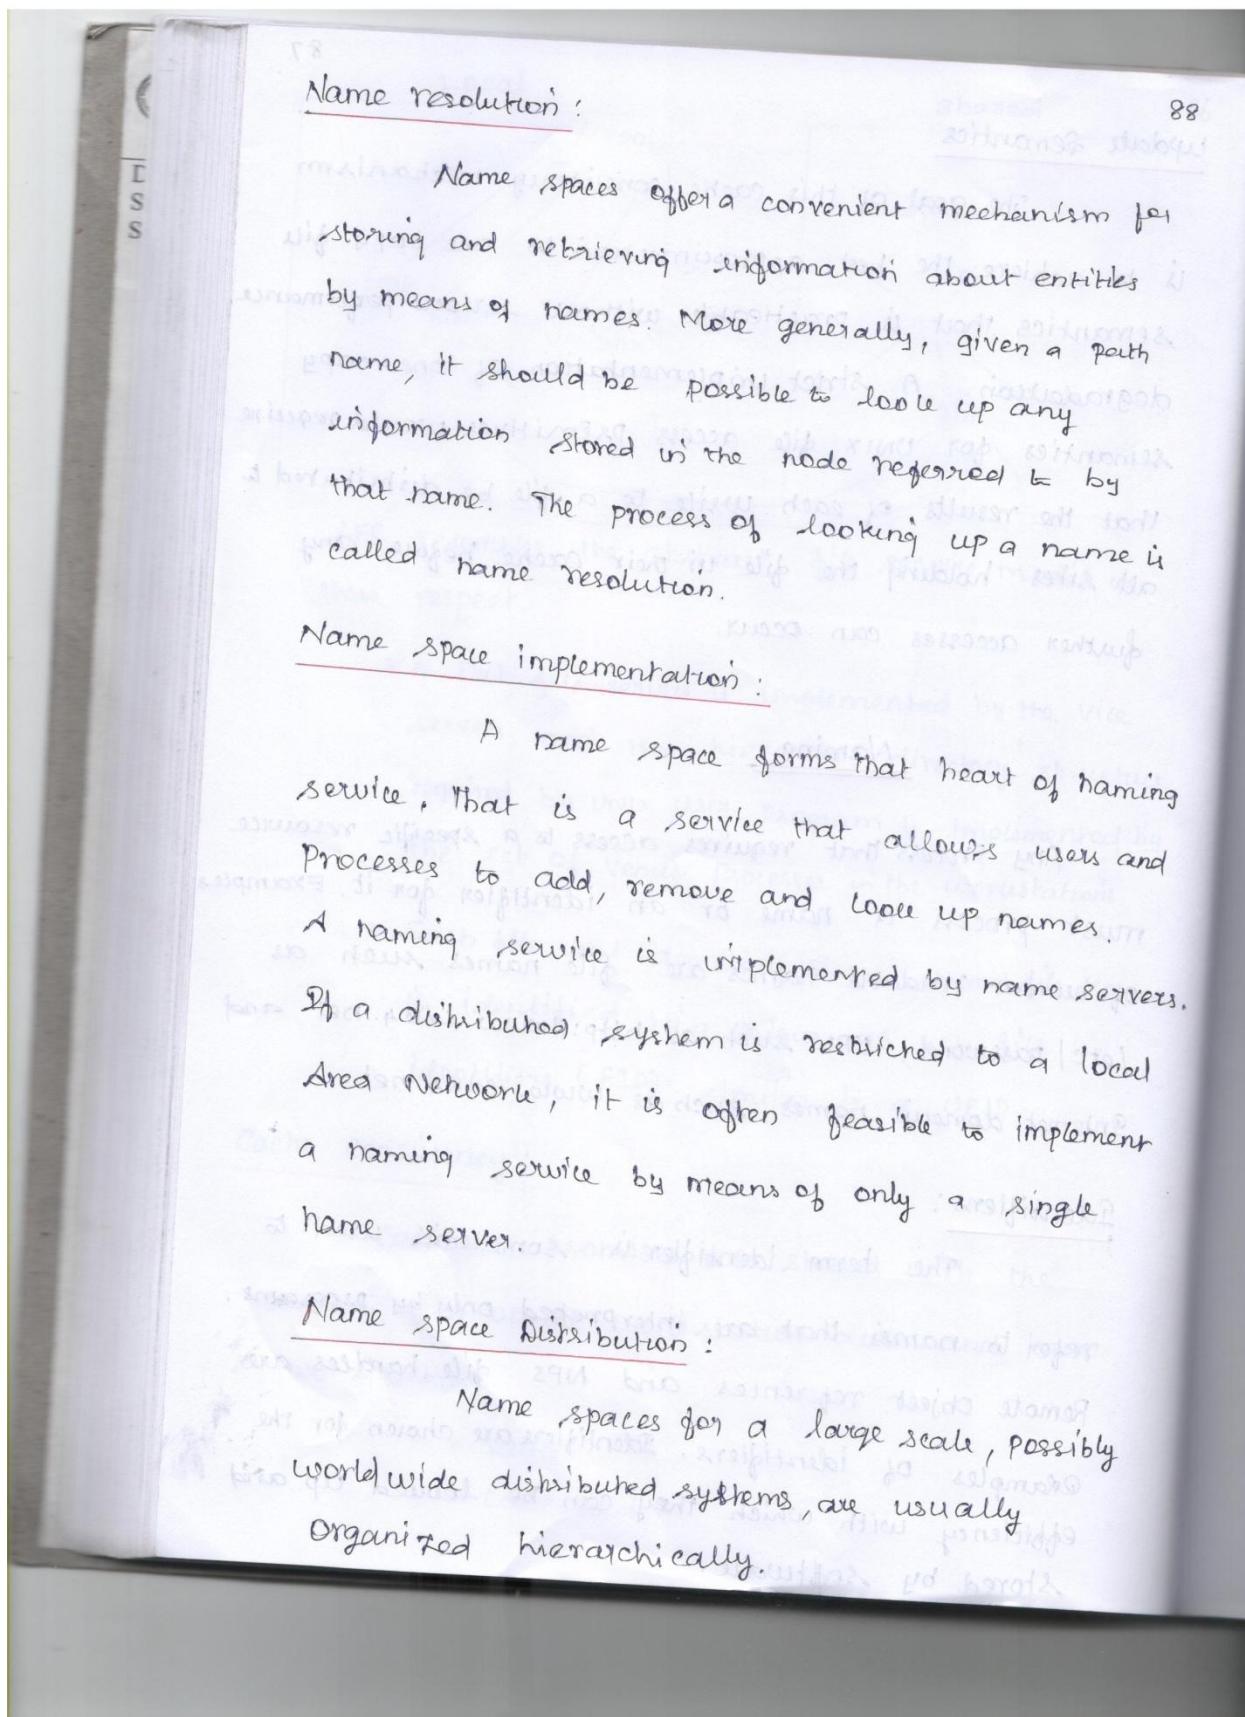
The goal of this cache consistency mechanism is to achieve the best approximation to one-copy file semantics that is practicable without serious performance degradation. A strict implementation of one-copy semantics for Unix file access primitives would require that the results of each write to a file be distributed to all sites holding the file in their cache before any further accesses can occur.

Naming

Any process that requires access to a specific resource must process a name or an identifier for it. Examples of human readable names are file names such as /etc/password, URLs such as http://www.cdu4.net and Internet domain names, such as www.cdu4.net.

Identifiers:

The term identifier is sometimes used to refer to names that are interpreted only by programs. Remote object references and NFS file handles are examples of identifiers. Identifiers are chosen for the efficiency with which they can be looked up and stored by software.



88

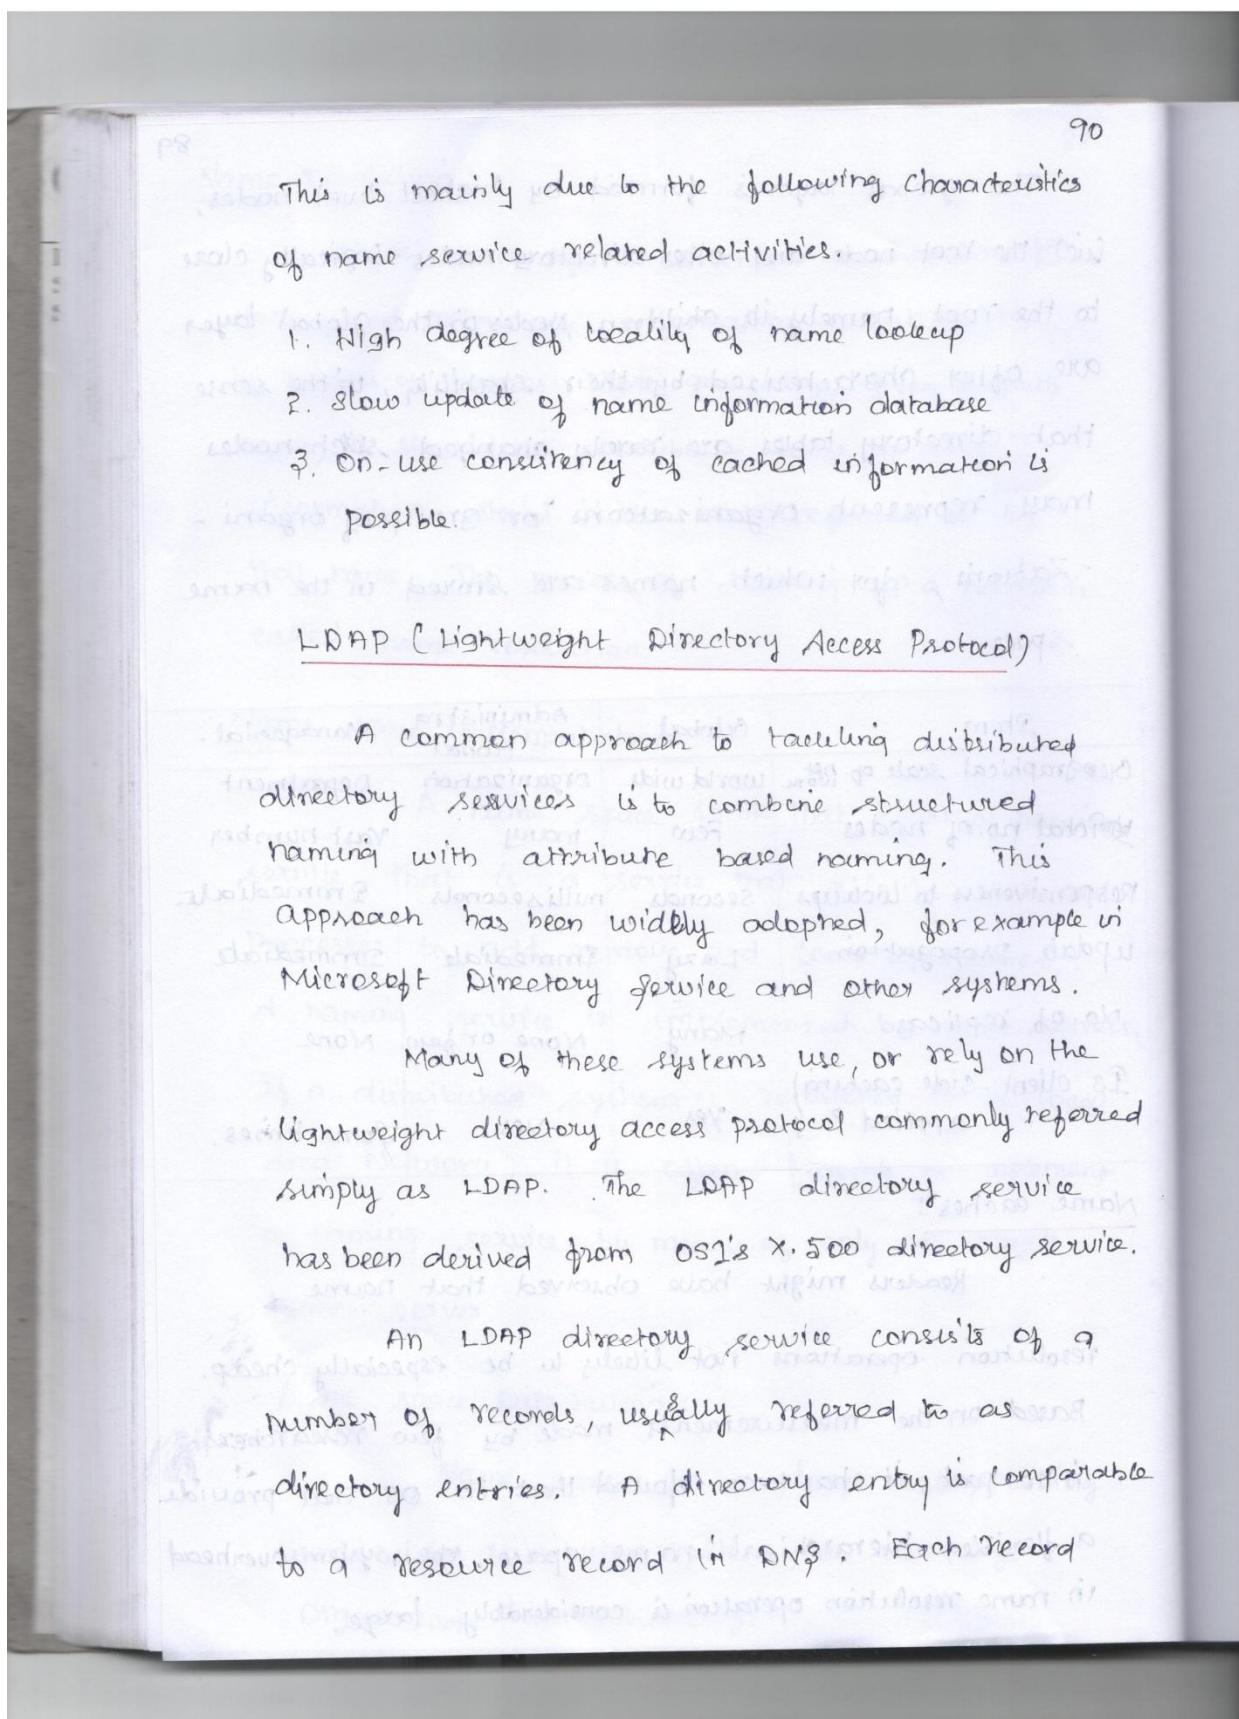
89

The global layer is formed by highest-level nodes, i.e. the root node and other directory nodes logically close to the root, namely its children. Nodes in the global layer are often characterized by their stability, in the sense that directory tables are rarely changed. Such nodes may represent organizations or group of organizations, for which names are stored in the name space.

| Item | Global | Administrational | Managerial |
|---------------------------------|------------|------------------|-------------|
| Geographical scale of network | World wide | Organization | Department |
| Total no. of nodes | Few | many | Vast number |
| Responsiveness to lookups | Seconds | milliseconds | Immediate |
| Update propagation | Lazy | Immediate | Immediate |
| No. of replicas | many | None or few | None |
| Is client-side caching applied? | Yes | Yes | Sometimes. |

Name caches:

Readers might have observed that name resolution operations not likely to be especially cheap. Based on the measurements made by few researchers in the past, it has been found that in OS that provide a flexible, hierarchical name space, the system overhead in name resolution operation is considerably large.



90

is made up of a collection of (attribute, value) pairs, where each attribute has an associated type. A distinction is made between single-valued attributes and multi-valued attributes.

91

As an example, a simple directory entry identifying the network addresses of some general servers, were as follows.

| Attribute | Abbr | Value |
|-------------------|------|--|
| country | C | NL |
| locality | L | Amsterdam |
| organization | O | Vrije Universiteit |
| Organization Unit | OU | comp.sc |
| Common Name | CN | Main server |
| Mail_Servers | - | 137.37.20.3, 130.37.24.6, 137.37.20.10 |
| FTP_Server | - | 130.37.20.20 |
| WWW_Server | - | 130.37.20.20 |

In our example, we have used a naming convention described in the LDAP standards, which applies to the first five attributes. The attributes organization and organization unit describe, respectively, the organization and the department associated with the data that are stored in the record.

